

## Chapter 1

# Distributed Software: Introduction

1. Definitions and Terminology
2. Developing distributed applications
3. Architecture
4. Design: middleware and services
5. Classes of distributed systems
6. Important architectures and platforms
7. Scalability and high availability
8. The Java family
9. The .NET family
10. This course

---

## Chapter 1

# Distributed Software: Introduction

## 1. Definitions and Terminology

1. Introduction
2. Examples
3. Definitions
4. Need for distributed systems
5. Why are distributed systems different ?

## 2. Developing distributed applications

## 3. Architecture

## 4. Design: middleware and services

## 5. Classes of distributed systems

...

## A loose definition

**distributed system =**

- collection of interacting processes
- appearing as single application to the end user

## “The network is the computer”

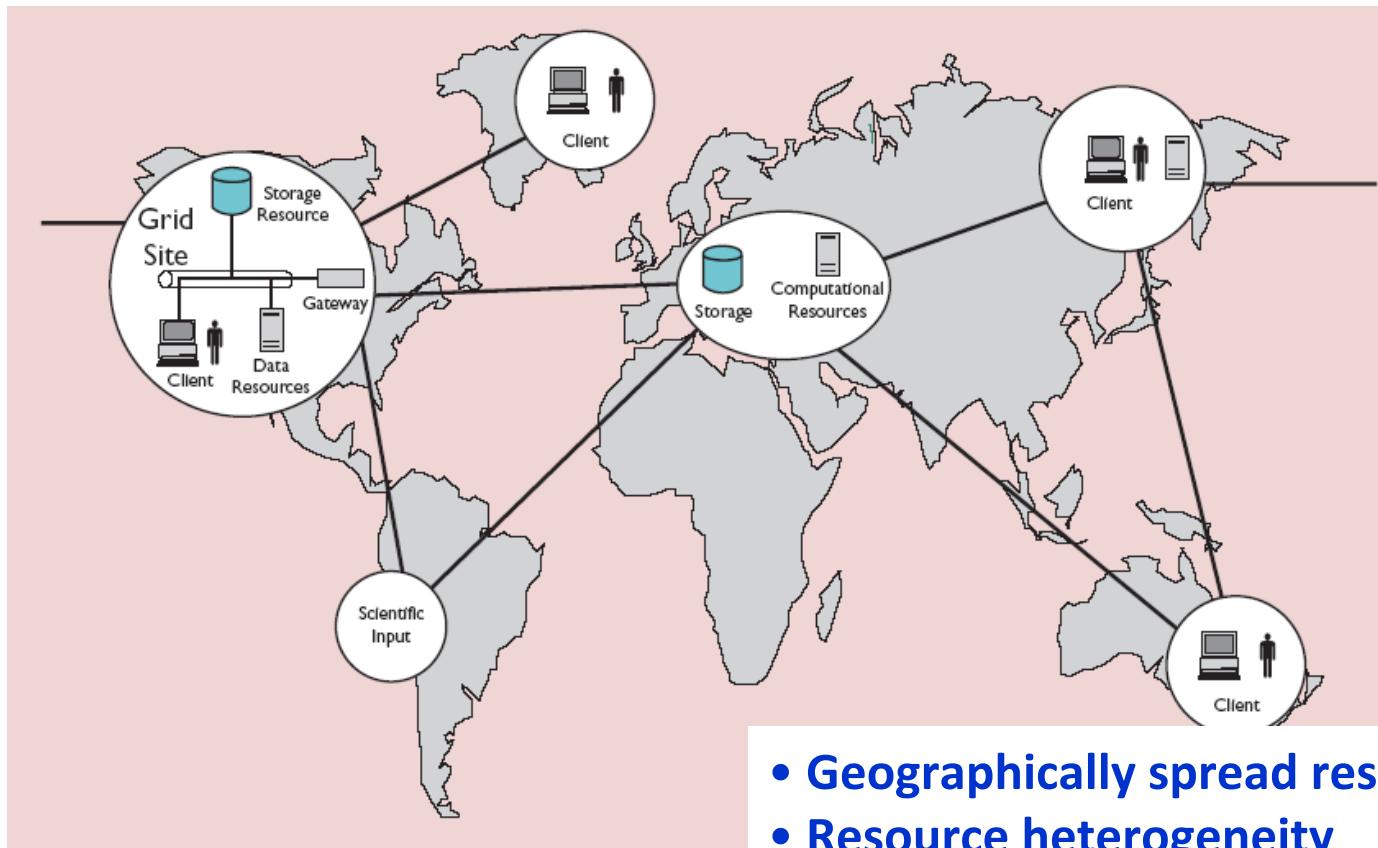
- search portals for huge information stores
- grid computing
- cloud management systems
- P2P applications
- communication tools
- online gaming
- Internet of Things applications

## Technology = standard programming tools + sockets ?

- specific, recurring problems due to distributed nature
- solve once, use many times
- middleware solves these problems in a generic way

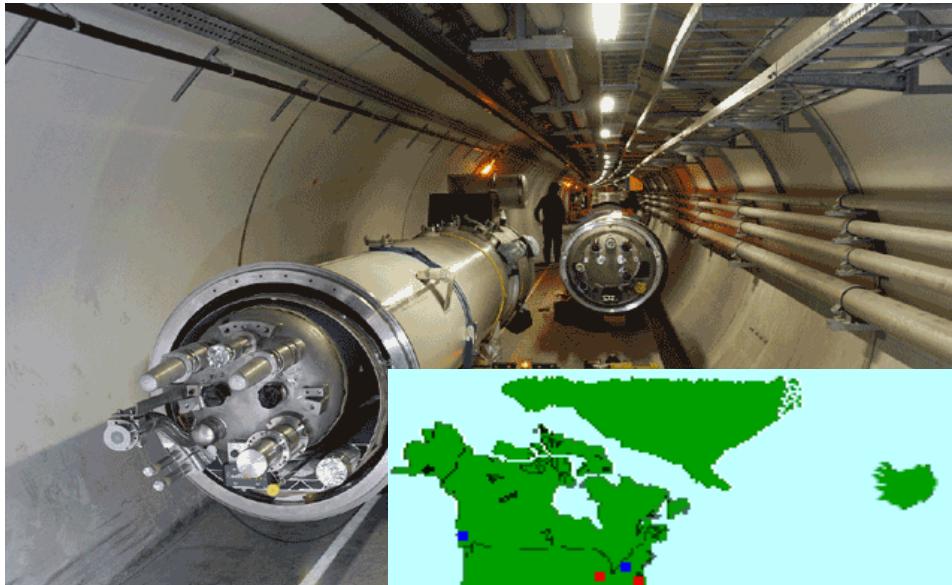
**“A computational grid is a hardware and software infrastructure that provides dependable, consistent, pervasive, and inexpensive access to high-end computational capabilities.”**

[Kesselman & Foster, 1998]



- Geographically spread resources
- Resource heterogeneity
- No single entity of control (no big boss)

## Large Hadron Collider (LHC) Computing Grid Project LCG

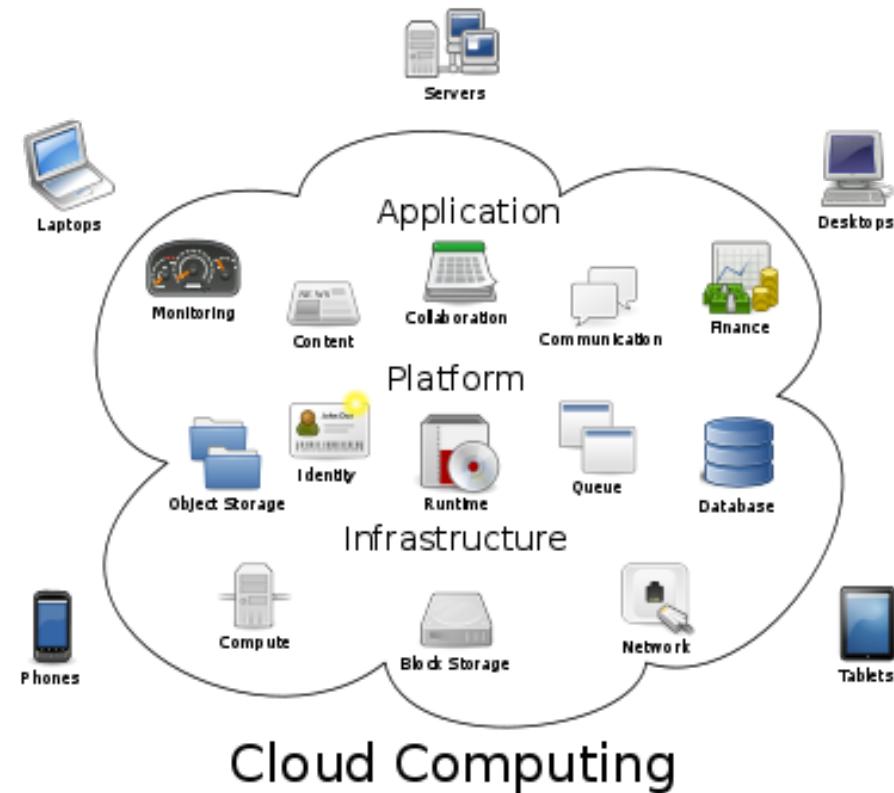


## Requirements

- > 6000 users
- 20 Pbyte/year [20  $10^6$  CDs]
- > 70 000 GFLOP

## Roughly defined

**Cloud computing = Grid computing  
+ virtualization  
+ elasticity  
+ business model**

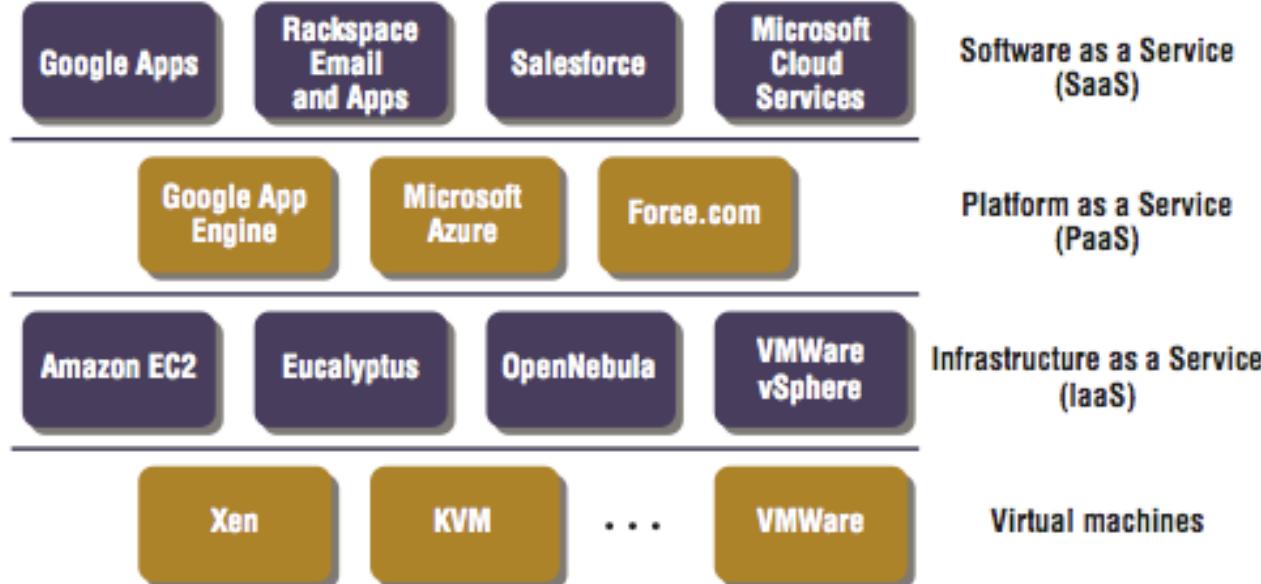


## The NIST Definition

“Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction. This cloud model is composed of five essential characteristics, **three service models, and four deployment models.**”

## Essential Characteristics

- on demand self-service
- broadband access
- resource pooling
- rapid elasticity
- measured service



## Service Models

- Software-as-a-Service (SaaS)
- Platform-as-a-Service (PaaS)
- Infrastructure-as-a-Service (IaaS)

## Deployment models

- Private Cloud
- Community Cloud
- Public Cloud
- Hybrid Cloud

## Some facts

Origin : sharing (often illegal) content (KazaA, BitTorrent, eMule, ...)

Huge user base

P2P internet traffic estimated at 70% of total traffic

## New applications

- Voice over IP (Skype)
- sharing computational power “desktop grids” (Folding@home)

## Benefits

- resources grow with number of users
- better scaling behaviour
- better robustness (no central servers)



## Nomadic scenarios

access guaranteed on other locations

(e.g. traveller accessing remote content, using local peripherals)

## Mobile scenarios

mobility (=moving user) support for applications

(e.g. taking location into account)

## Ubiquitous computing scenarios

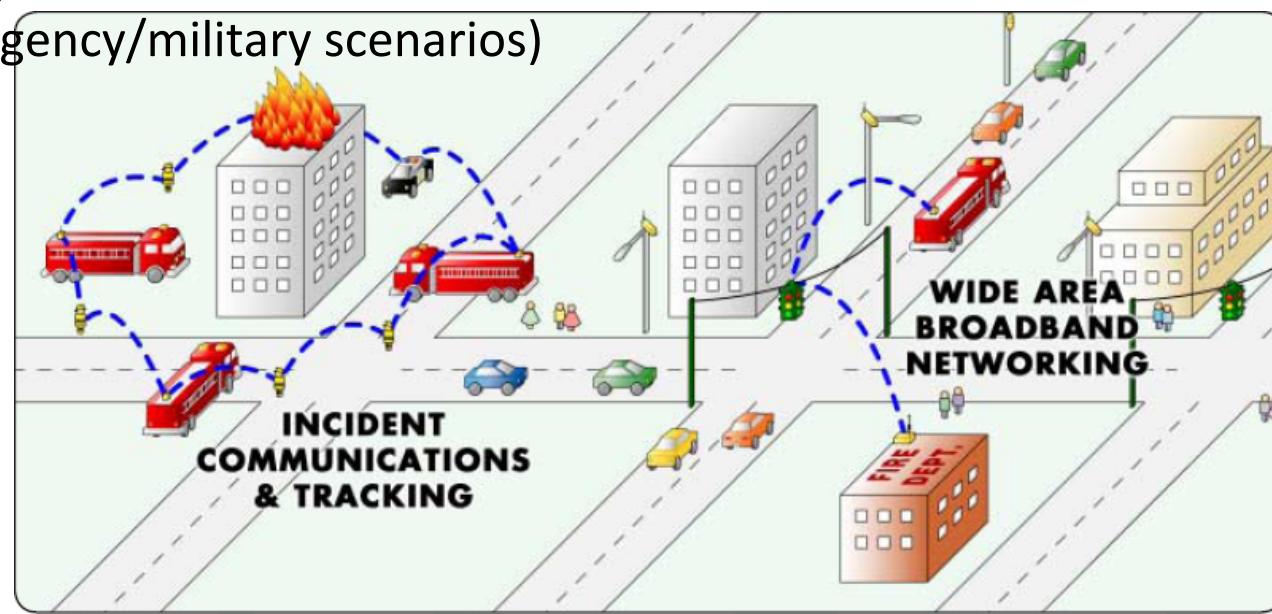
application fed with sensor info, steers actuators

(e.g. cleaning robot)

## Ad-hoc scenarios

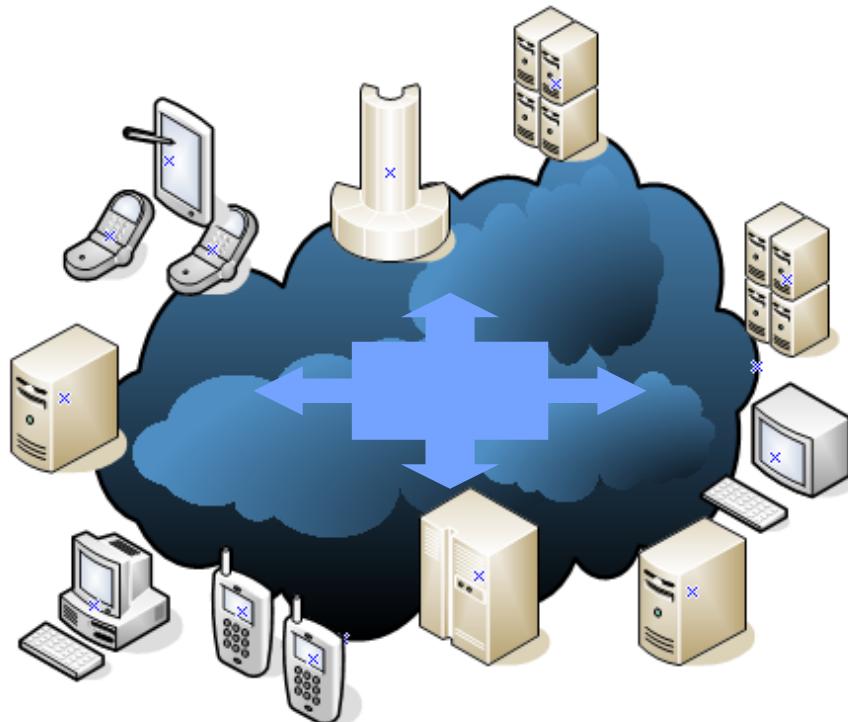
rapid deployment of infrastructure

(often emergency/military scenarios)



**A distributed system** = a system where

- hardware and software components are located at **networked** computers
- components communicate and coordinate their actions **ONLY by passing messages**



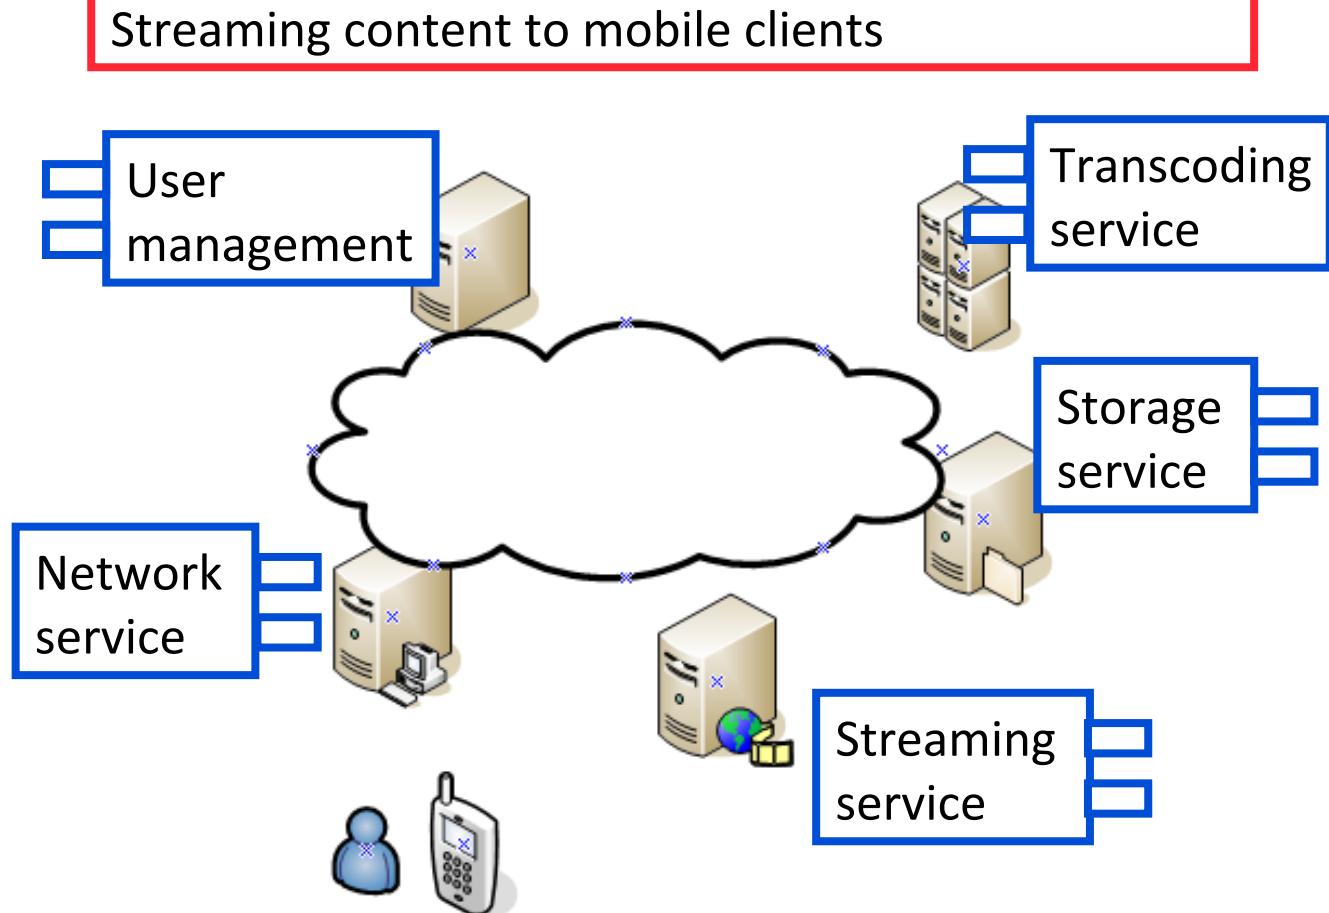
## Consequences

- no limit on spatial extent
- no global time notion
- almost always concurrent execution
- (partial) failures likely to happen

A **service** = part of a computer system managing a collection of related resources, presenting their functionality to users and applications

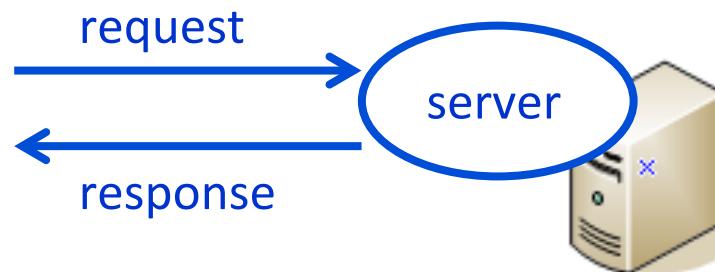
## Examples

- storage service
- print service
- streaming service
- connectivity service
- transcoding service

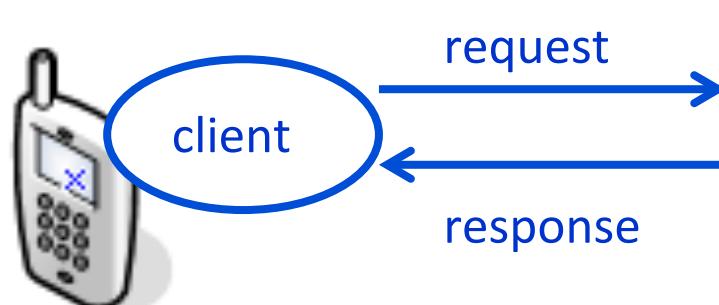


**A server** = running **process** on networked computer

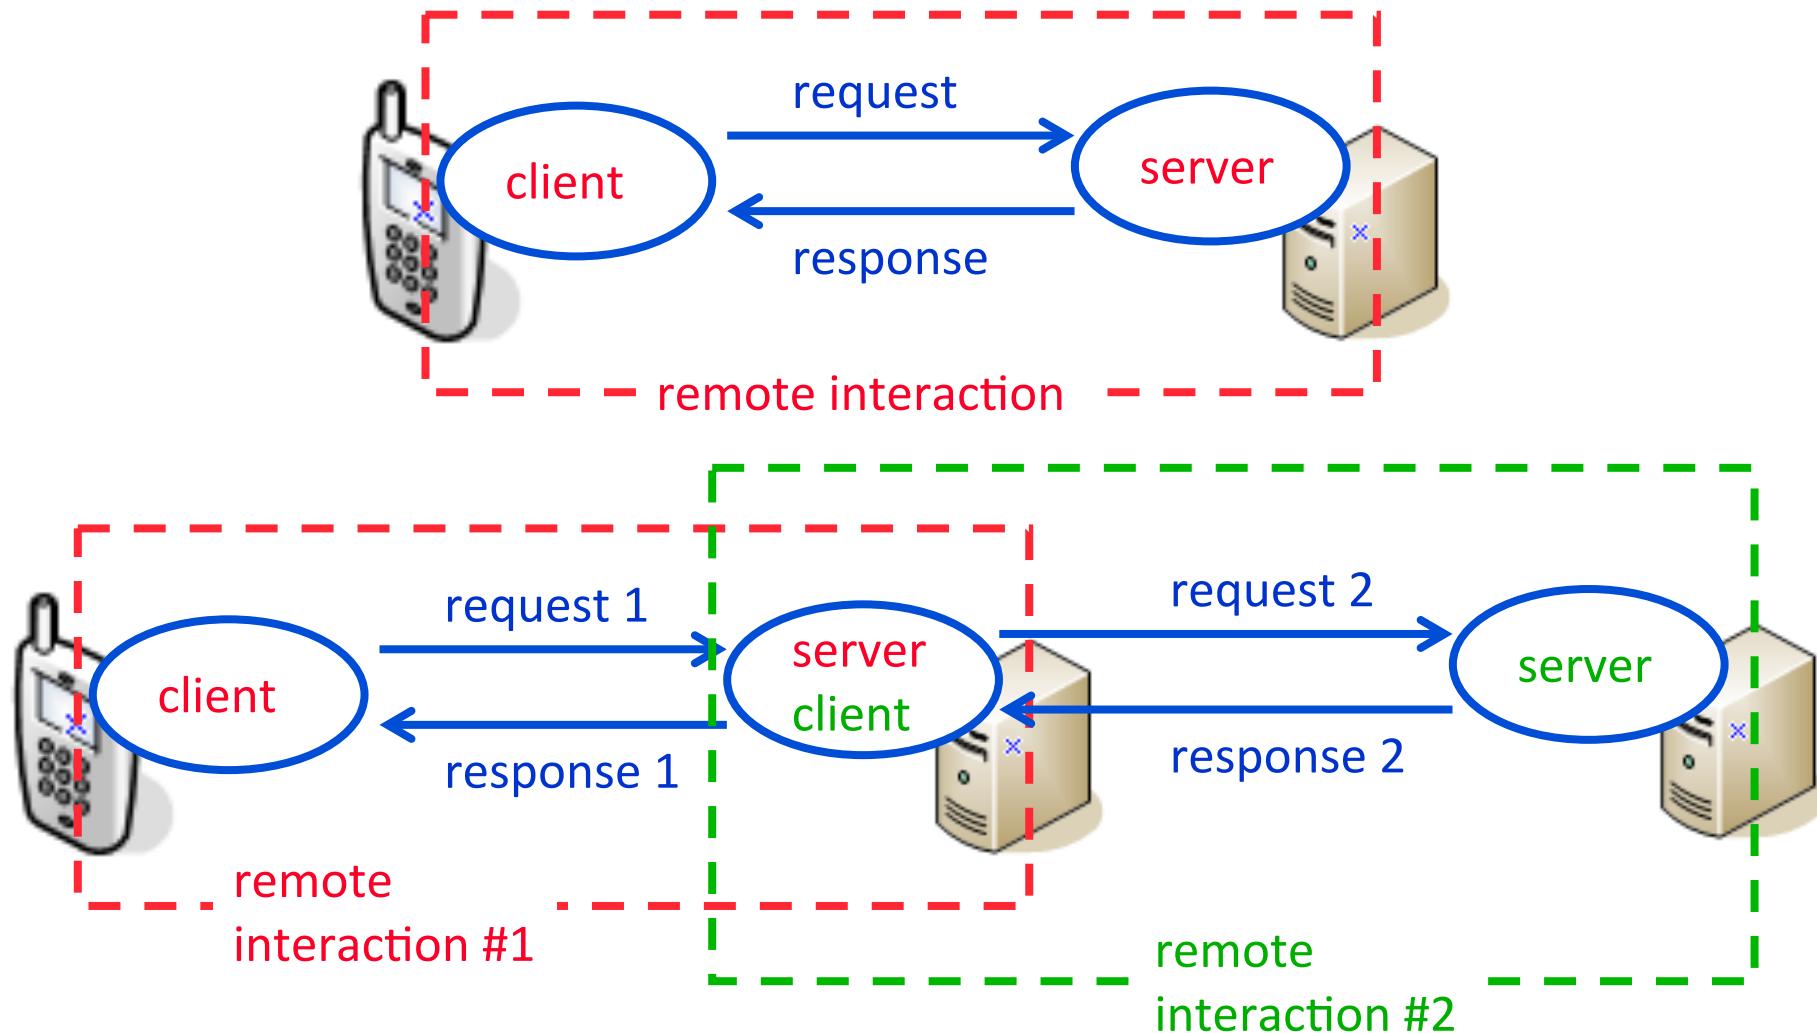
- accepting requests to perform a service
- responding appropriately



**A client** = running **process** on networked computer **sending** service requests to servers



A remote invocation = complete interaction between client and server to process a single request



## Many problems

- no limit on spatial extent, difficult to manage
- no global time notion
- almost always concurrent execution
- (partial) failures likely to happen

## Many advantages

- resource sharing
  - information resources
  - hardware resources
- scalability
  - “easy” to adapt to larger user base
- fault tolerance
  - “easy” to cope with failures



# Distributed system vs. standalone system

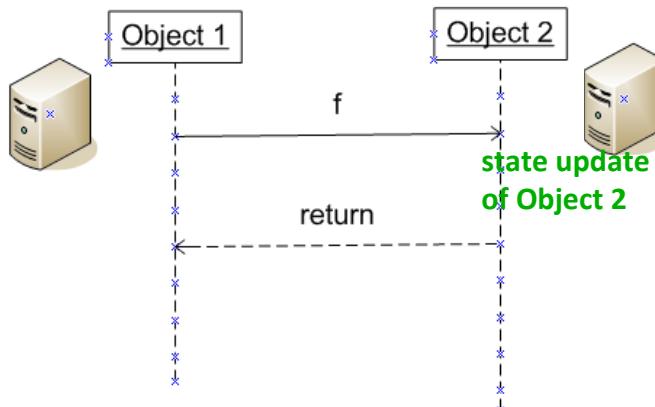
1. Definitions and terminology
5. Why different ?

Standalone system	Distributed system
Communication between processes possible through shared memory.	Communication between processes only through the network.
Global state possible through shared memory.	No global state.
Local operating system can be used to get a shared time value.	No global notion of time.
Local operating system offers primitives to safely share resources.	Synchronisation between processes through network communication only.
Failing processes easy to detect (communication is reliable).	(Partial) failures difficult to detect (because of unreliable communication).
Infrastructure known beforehand.	Infrastructure can vary dynamically (new servers/nodes can become available).
Components are located on the standalone machine.	Location of components can vary dynamically.
Local operating system enforces security.	Security threat due to distributed nature (possibly vulnerable communication link).

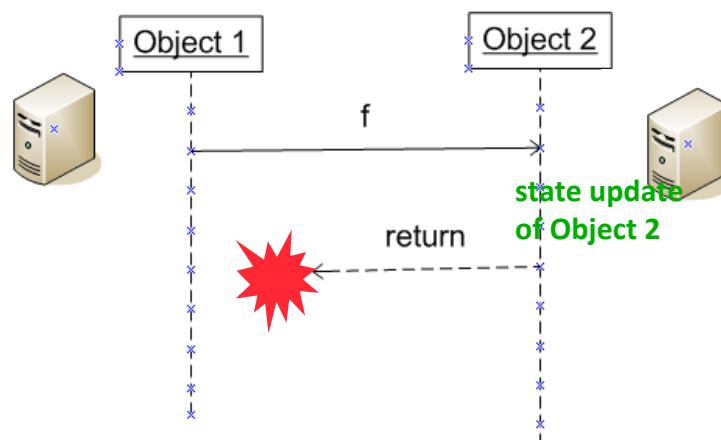
# Faulty communication or faulty process ?

1. Definitions and terminology
5. Why different ?

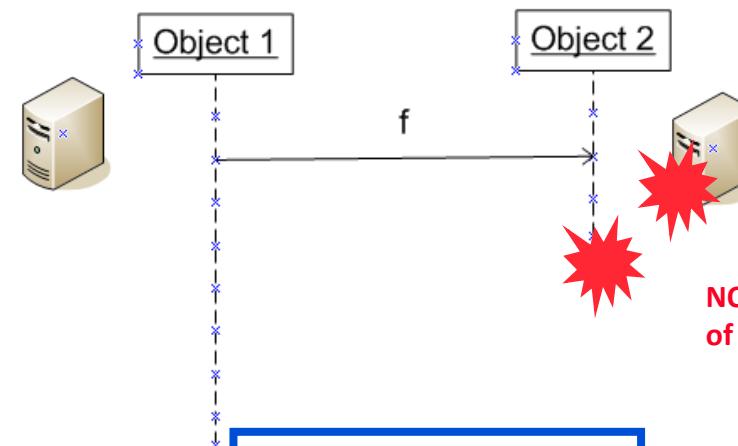
Normal scenario



Failure scenario's



communcation  
failure [1]



process  
failure [2]

NO state update  
of Object 2

How can Object 1 distinguish between [1] and [2] ?

## **Chapter 1**

# **Introduction**

- 1. Definitions and Terminology**
- 2. Developing distributed applications**
  1. The development process
  2. System models
  3. Challenges for developing distributed app's
- 3. Architecture**
- 4. Design: middleware and services**
- 5. Classes of distributed systems**
- 6. Important architectures and platforms**

...

## Requirements analysis phase

- required functions
- non-functional constraints

## Architectural phase

Formalised requirements  
UML-diagrams capturing requirements

## Design phase

High level architecture  
subsystems and interfaces  
Detailed design  
(in parallel for each subsystem)

WHAT ?  
[OOA]

## Implementation phase

Code + documentation (in parallel)

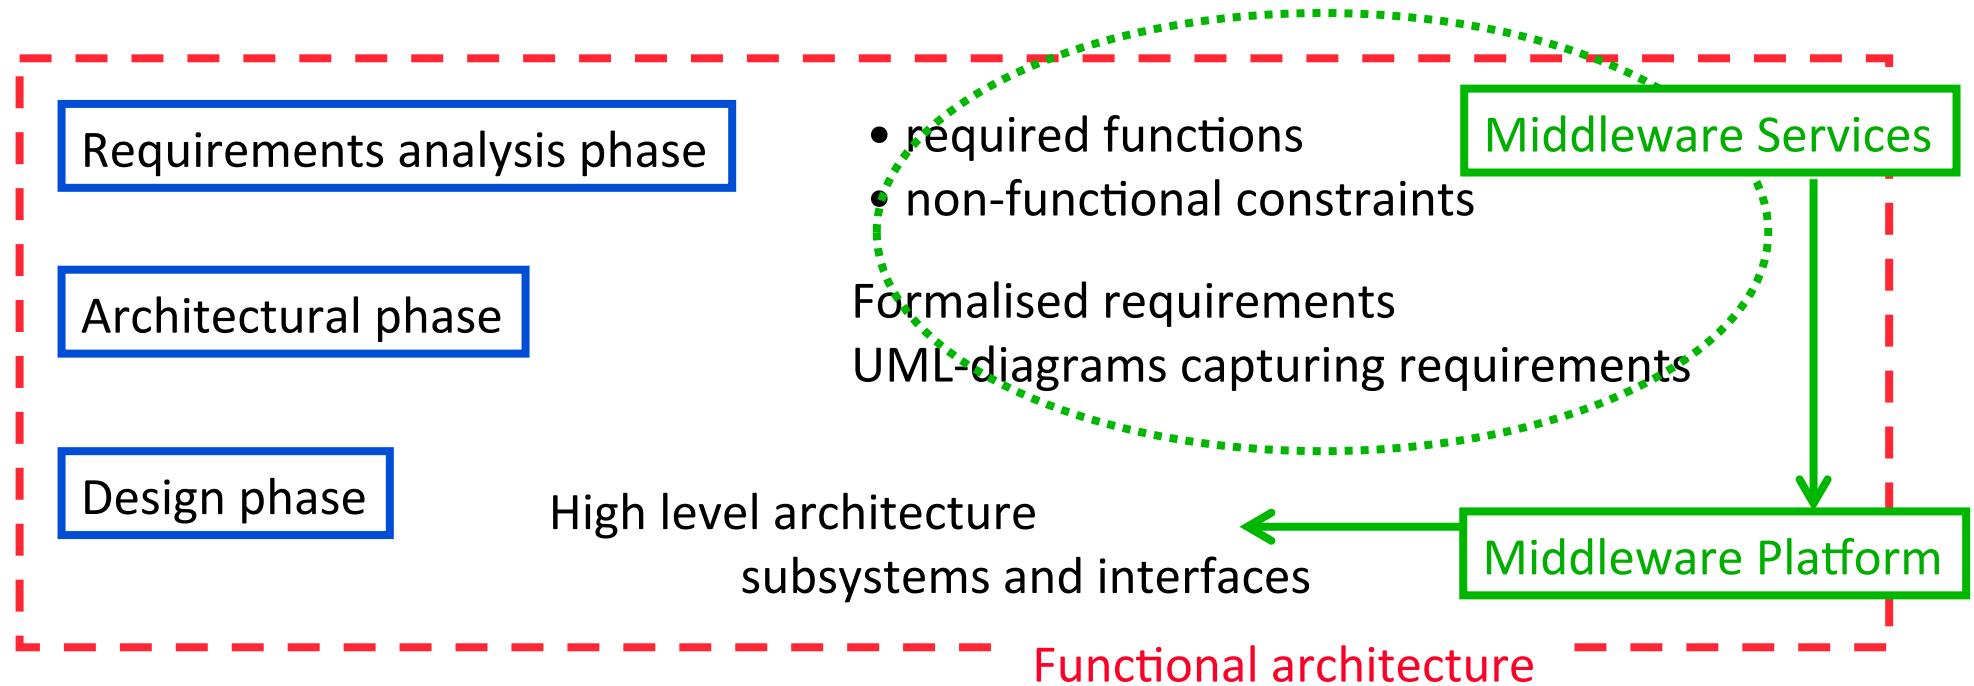
HOW ?

## Integration phase

Combine realised subsystem code

# Developing distributed applications

## 2. Developing distributed app's 1. Development process



### Logical architecture

how will the subsystems interact  
(e.g. eventing, through remote method call, data base, ...)

### System architecture

client-server or P2P ?

### System models

non-functional constraints relating to distributed nature

## System models

-> capture non-functionals of the requirements related to distributed nature

### -> Interaction model

can we give time guarantees on network communication/process execution ?

### -> Failure model

which failures should the application cope with ?  
how will we detect failures ?  
what will we do in case of failures ?

### -> Security model

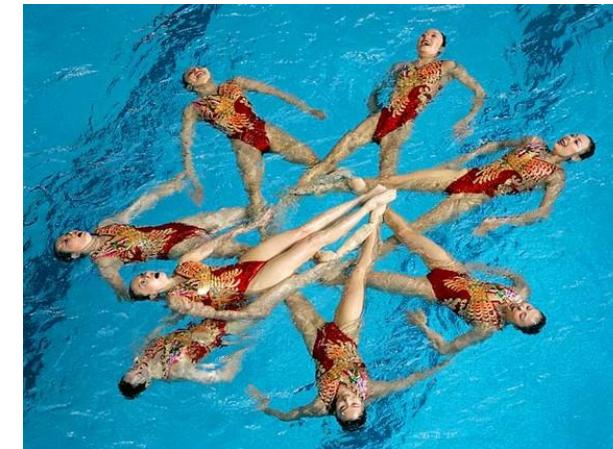
who can do what in the system ?



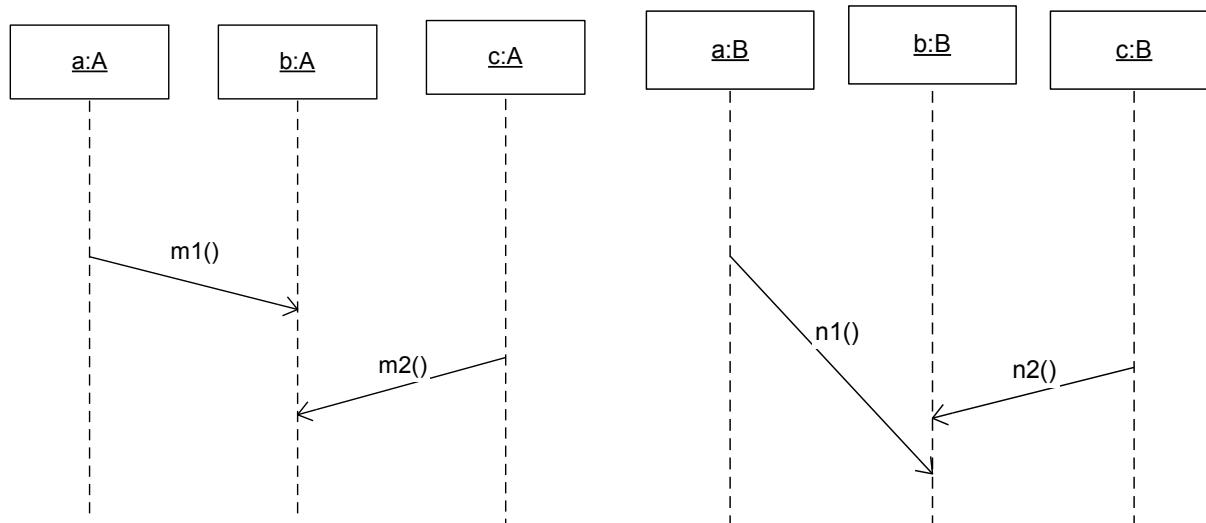
A **synchronous system** is a system where

1. The time to execute each process step has a known lower and upper bound
2. Each message is received within a known bounded time
3. Each process controlled by local clock with known bound for drift rate w.r.t. real time

An **asynchronous system** is a system where one of the above mentioned conditions does not hold.



Asynchronous system  
needs special algorithms  
to detect event order



## Omission failures

a component fails to perform an action

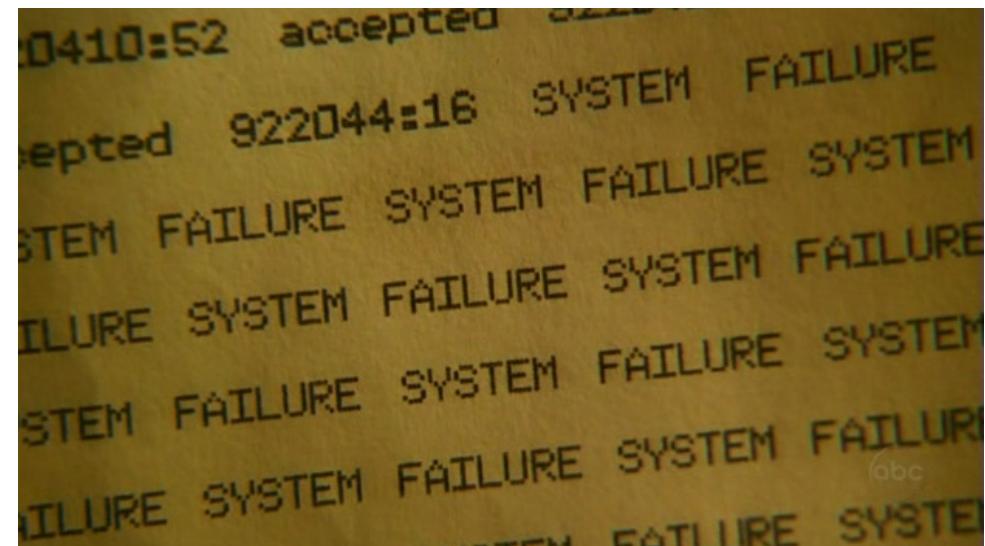
- Process omission : a process does not fulfill a request
- Channel omission: a communication link fails to deliver a message

## Byzantine failures

arbitrary behavior (e.g. a process giving the wrong result)

## Timing failures

a synchronous system is violating one (or more) of its timing constraints



## Detection mechanisms

heartbeat

timeout

## Security threats

### Process

server : enemy can use intercepted credentials

client : enemy can pretend to be server

### Channels

copy, change, replay, create messages

### Privacy issues

## Solution

Authentication/Authorization infrastructure

+ Cryptographic solutions

(shared secret, asymmetric keys, ...)



## Challenge

### Heterogeneity

physical/L2 network  
hardware  
operating system  
programming languages

## Solution

Common (standard) protocols  
Virtualization  
Virtual machines  
Middleware

## Security

privacy (e.g. financial/medical data)  
secure accounting  
prove identity  
viruses  
(distributed) denial of service attacks  
mobile code

cryptographic infrastructure  
authorization/authentication  
virus scanning  
firewall, VPN  
sandboxing

## Scalability

scaling of infrastructure cost  
    physical cost =  $O(\#users)$   
performance scaling  
scaling of logical resources  
performance bottlenecks

clever design  
automatic replication of components

# Challenges for developing distributed app's

2. Developing distributed app's
  3. Challenges
- 

## Challenge

### Failures

failure detection  
failure handling

high availability

## Concurrency

resource access conflicts

## Solution

polling, heart beat, checksum,...  
masking : retransmit messages,  
fail over  
tolerate : inform user of  
(partial) failure  
recover : roll back to ensure  
consistency  
redundancy and replication

lock mechanisms  
synchronization primitives

# Challenges for developing distributed app's

2. Developing distributed app's
3. Challenges

## Challenge

### Transparency

#### Access transparency

identical operations to access local and remote resources

#### Location transparency

access of resources without location knowledge

#### Concurrency transparency

provide safe resource sharing amongst concurrent processes

#### Replication transparency

replicated resources are used implicitly

#### Failure transparency

hide hardware/software failure

#### Mobility transparency

resources/clients can move without affecting system operation

#### Performance transparency

system can be reconfigured according to load conditions

#### Scaling transparency

system can expand without changing structure and algorithms

## Solution

### Middleware



---

## **Chapter 1**

# **Distributed Software: Introduction**

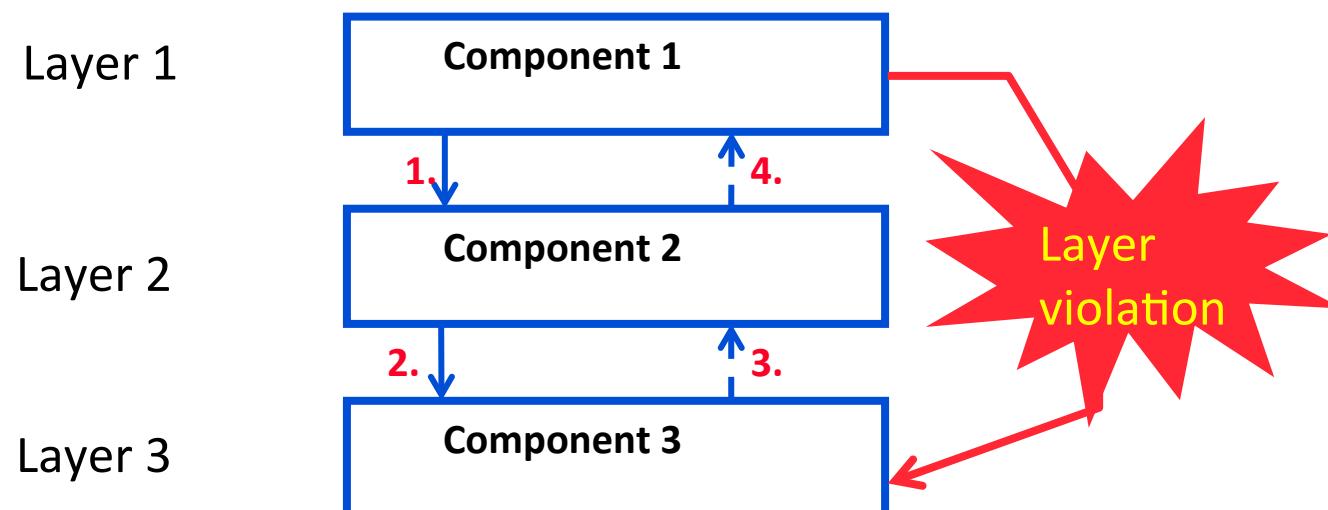
- 1. Definitions and Terminology**
- 2. Developing distributed applications**
- 3. Architecture**
  - 1. Logical architecture**
  - 2. System architecture**
- 4. Design: middleware and services**
- 5. Classes of distributed systems**
- 6. Important architectures and platforms**

...

= architecture capturing how subsystems will interact  
also referred to as “architectural style”

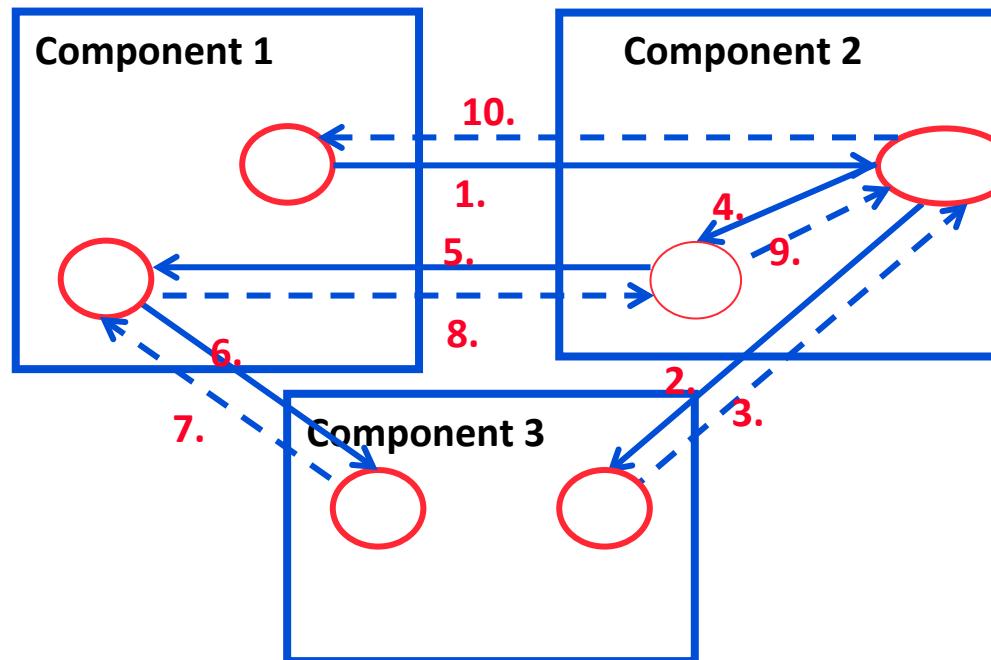
## Layered architecture

- layer only interacts with neighbour
  - + reduced number of interfaces, dependencies
  - + easy replacement of a layer
  - possible duplication of functionality



## Interacting objects

- no predefined interaction pattern
- + highly flexible
- complex to manage and maintain



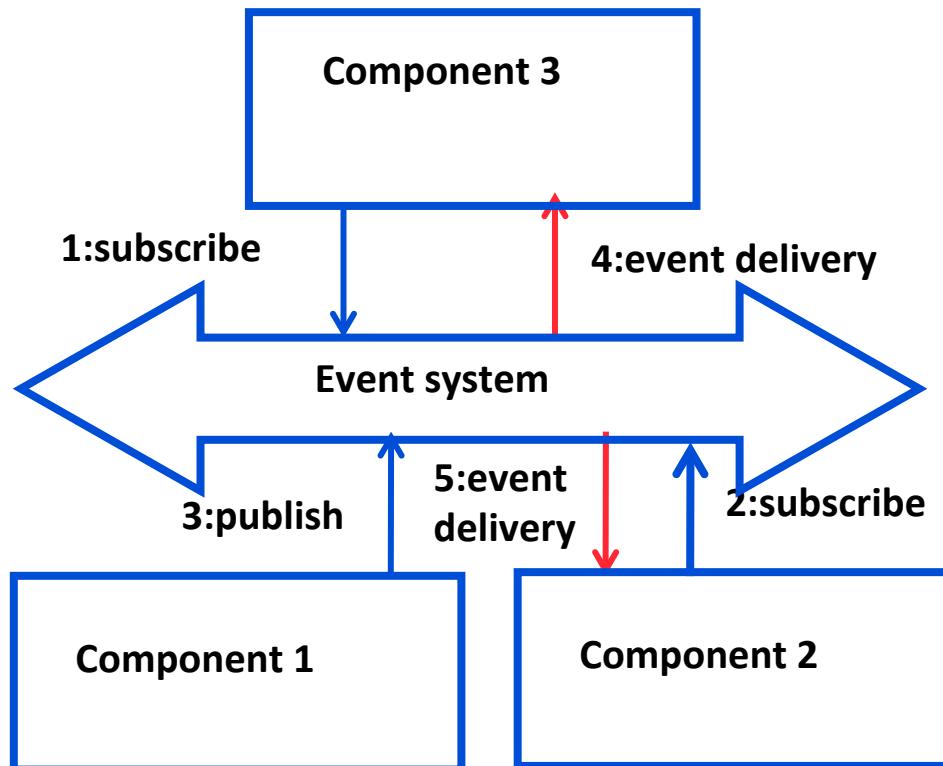
## Event based interaction

“publish-subscribe” style

+ loose coupling of components

related: message based interaction (also decoupling in time)

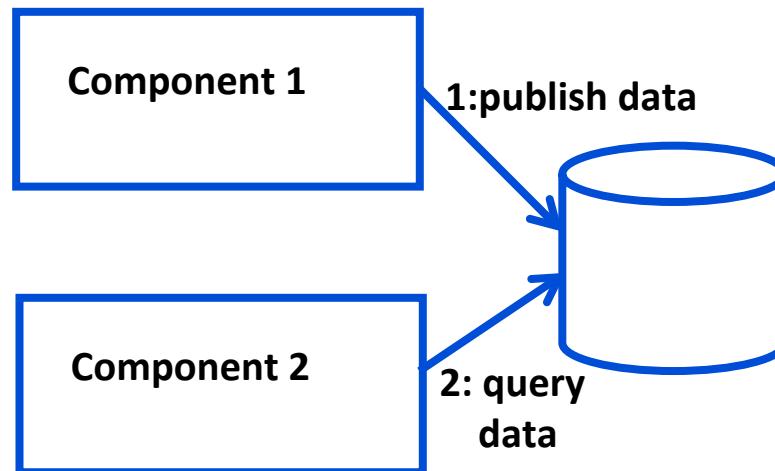
often used to integrate legacy systems



## Data centric architecture

only interaction through shared data base

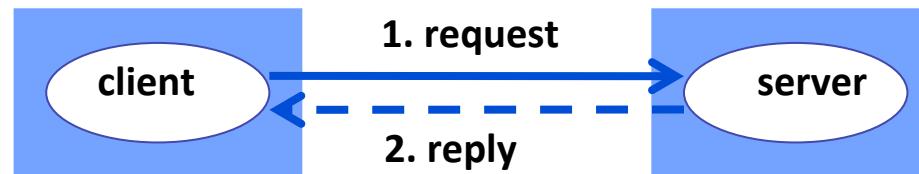
- + loose coupling of components
- possibly slow (central bottleneck, locking, ...)



How to map logical components to actually deployed components  
(replicated ?, P2P, pure client server, ...)

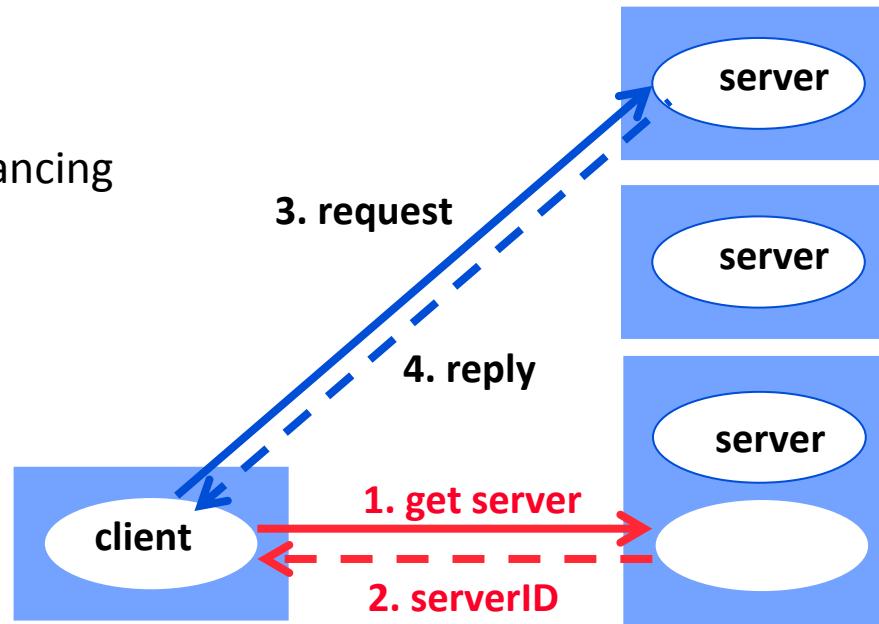
## Client-server architectures

“simple” client-server

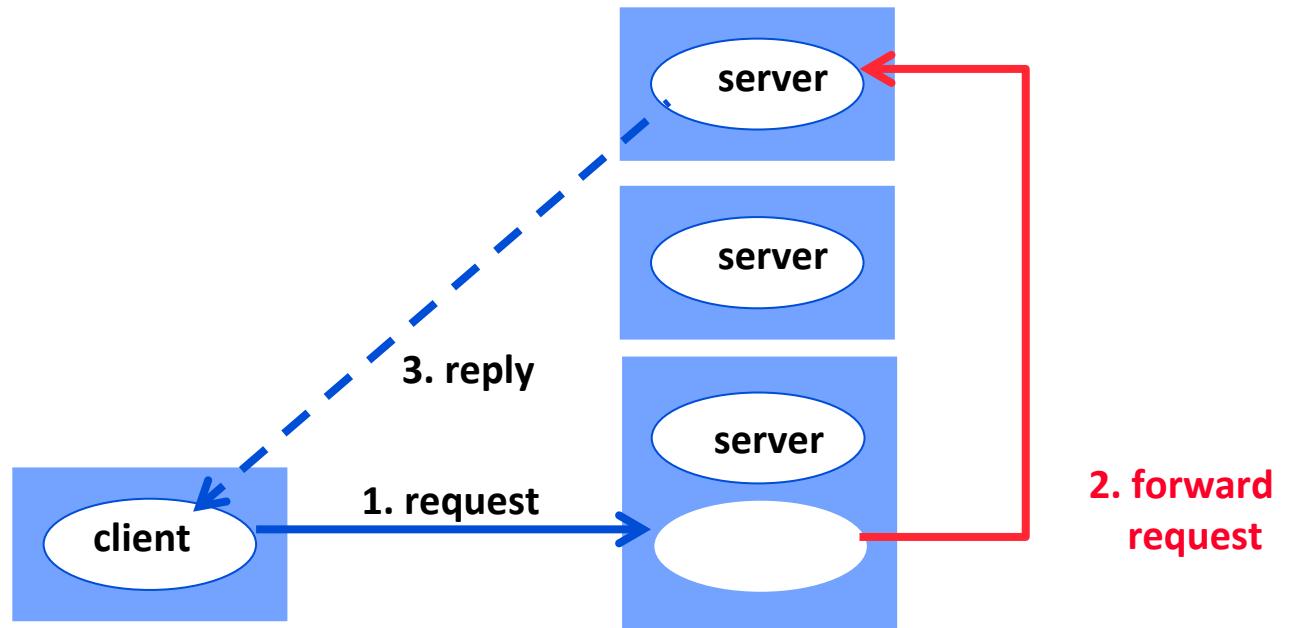


client- multiserver  
(explicit server lookup)

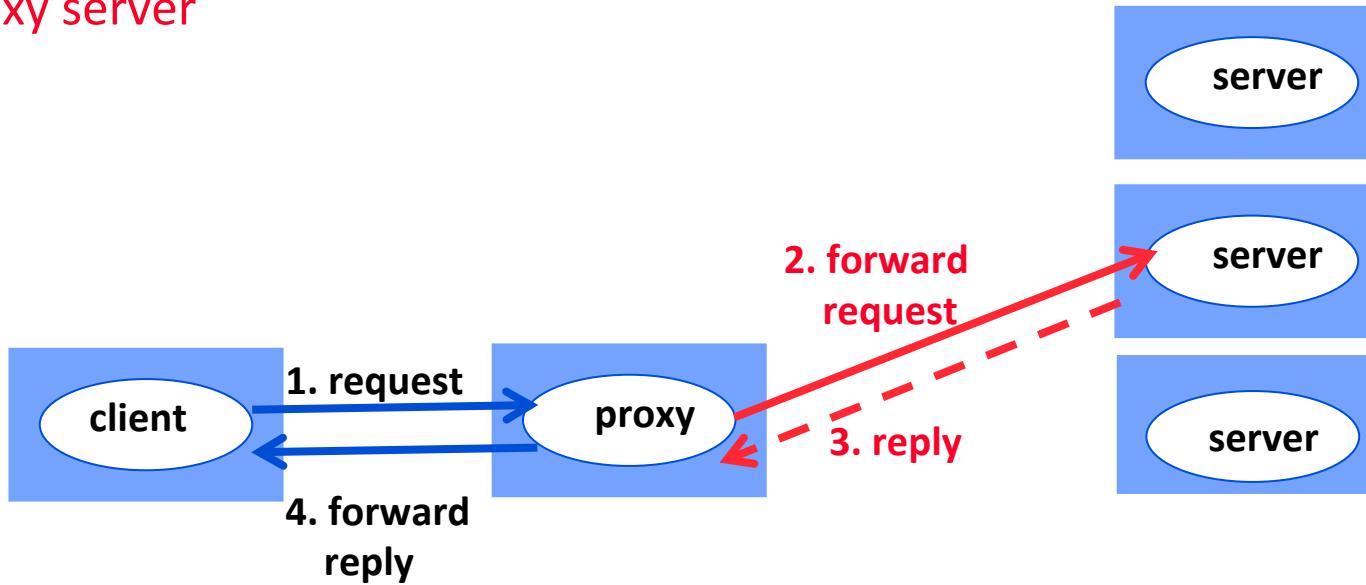
e.g. DNS based load balancing



client- multiserver  
(implicit server lookup)



proxy server



**servent** = process issuing requests + fulfilling requests

peering components use logical network (overlay) to interact

### Unstructured P2P :

- links have no other meaning than communication
- unstructured search (e.g. flooding) : resource consuming

### Structured :

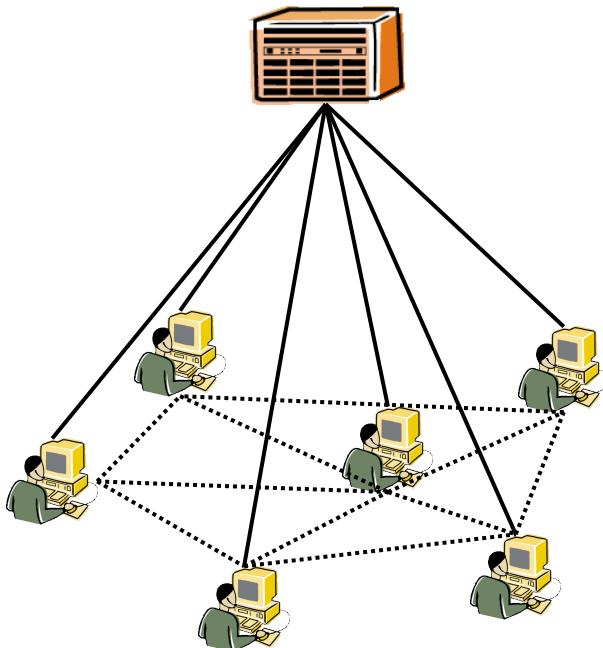
- logical link has relation to service offered
  - structured search much more efficient
- important data structure : Distributed Hash Table (DHT)

## Why P2P ?

- scalability (10000 – 100000 users not exceptional)
- gradual scaling (“organic growth”)
- robustness

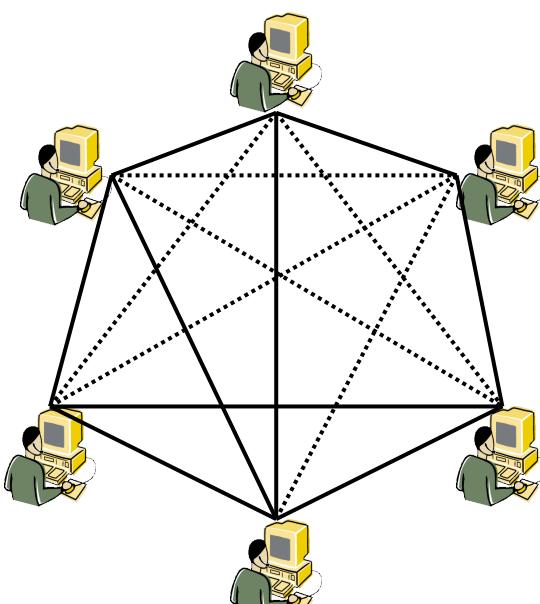
## Unstructured

### Mediated P2P



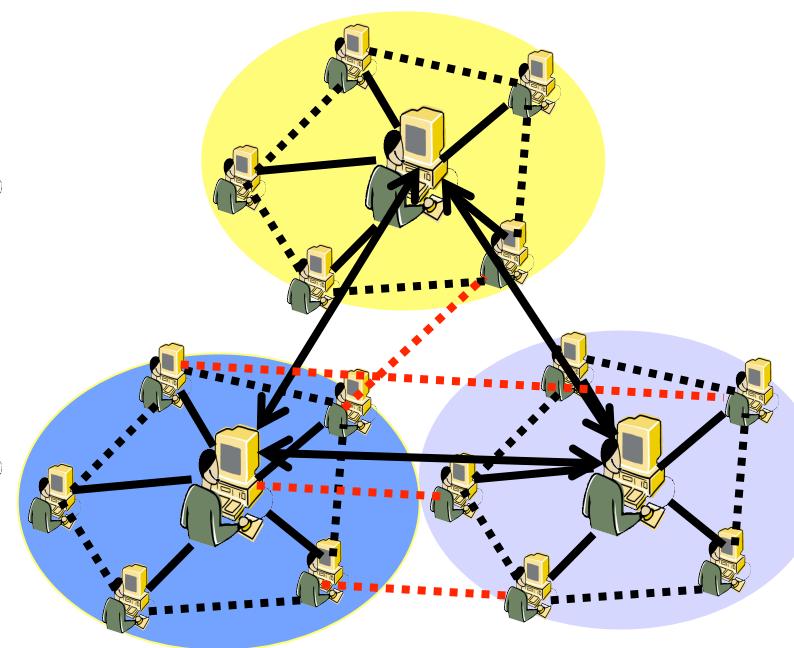
Napster  
Audiogalaxy

### Pure P2P



Early Gnutella  
FreeNet

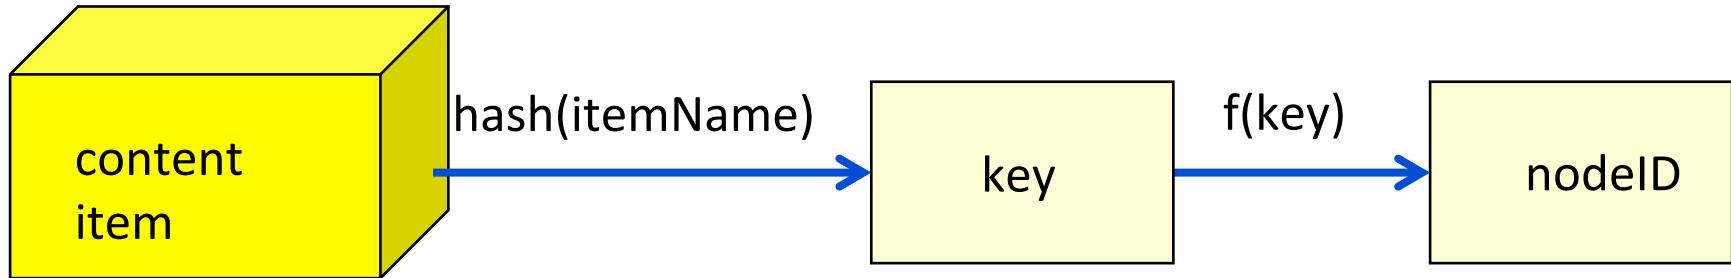
### Hybrid P2P



KazaA  
control traffic (index, lookup)  
data traffic (download, upload)



## Structured

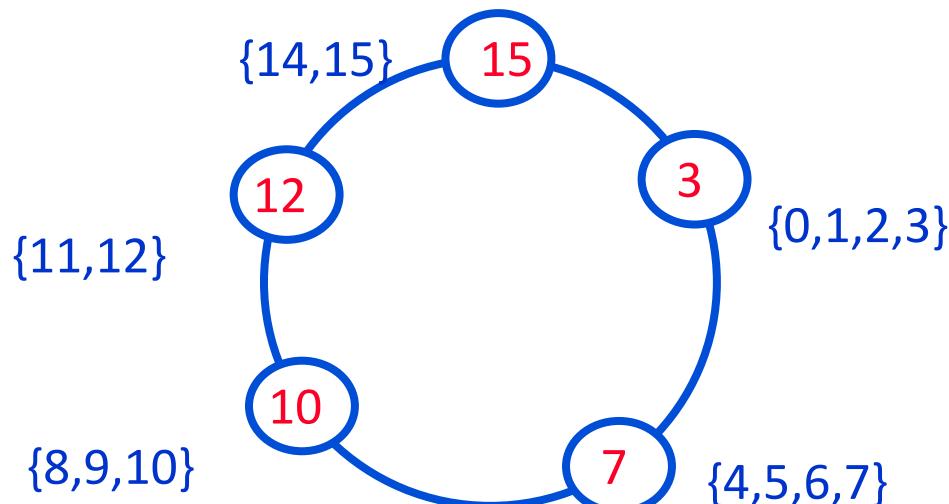


## Chord

each node has integer nodeID

e.g. consider 5 node network  $\{\text{nodeID}\} = \{3, 7, 10, 12, 15\}$

$\{\text{key}\} = \{0 .. 15\}$



see chapter “P2P” for more details

## Chapter 1

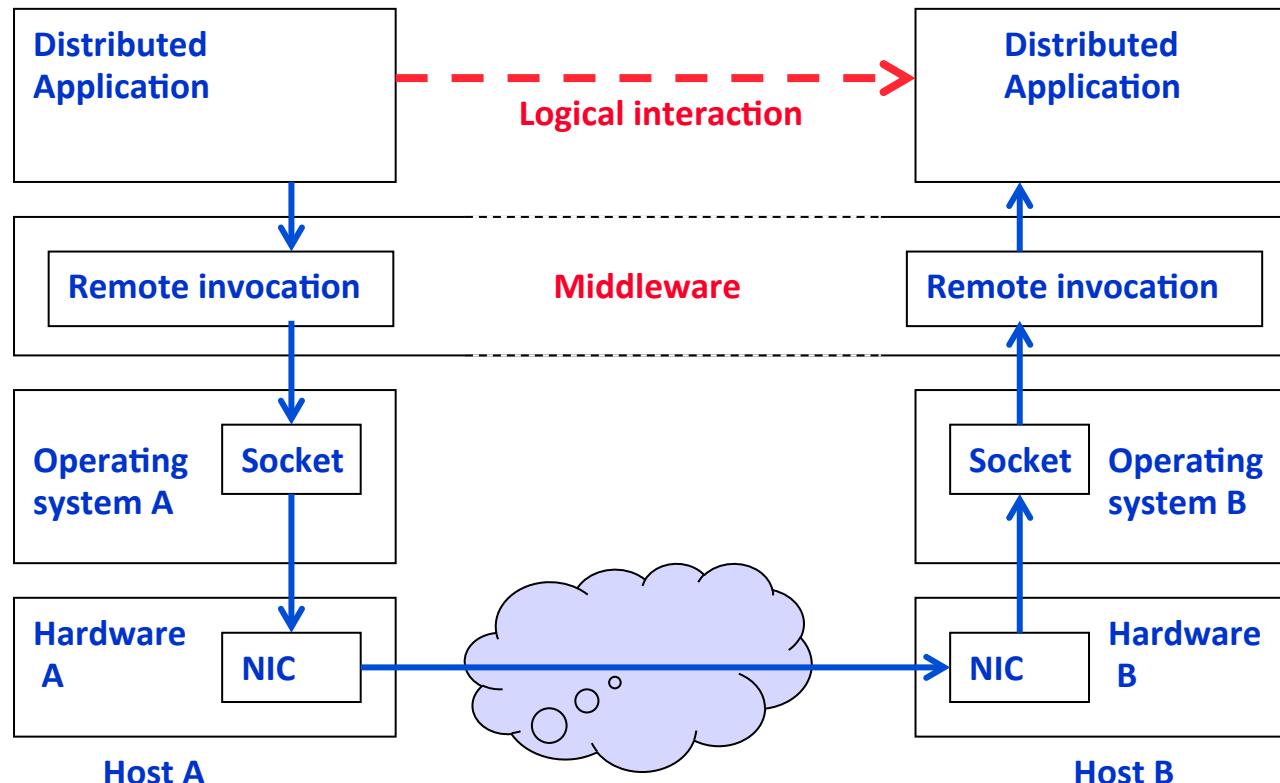
# Distributed Software: Introduction

1. Definitions and Terminology
2. Developing distributed applications
3. Architecture
- 4. Design: middleware and services**
5. Classes of distributed systems
6. Important architectures and platforms
- ...

## 1) Abstract hardware/software platform

- operating system
- implementation language
- ...

## 2) Realize transparency – focus on location transparency



### 3) Provide generic services

#### Naming service

associate logical names to remote entities

#### Discovery and registration service

easy service lookup and registration

#### Transaction service

support distributed (database) transactions

#### Security service

support to authenticate and authorize users

#### Event service

distributed notification of events

#### Data replication

offer the same data at different locations to optimize data access

#### Persistence service

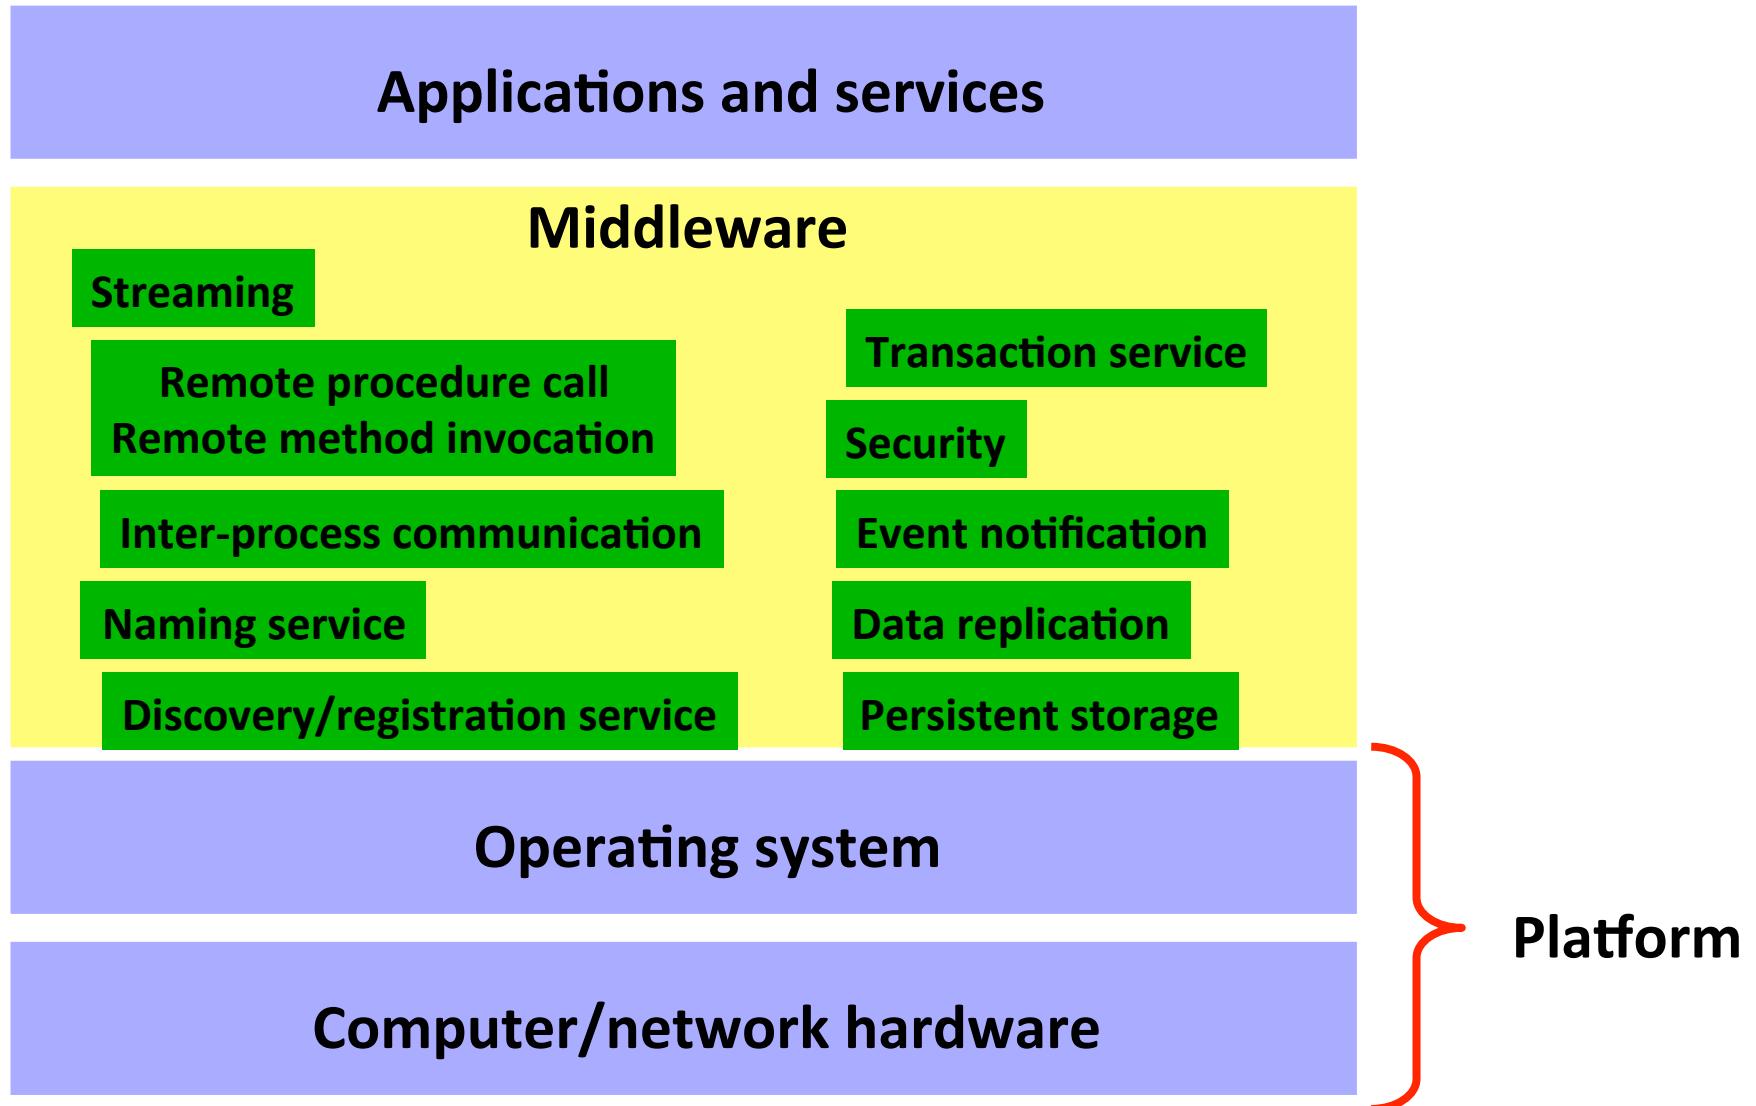
to transparently offer data persistence

(i.e. data is automatically stored in a persistent medium)

#### Life cycle service

managing service components

(start up of services, shut down, etc.)



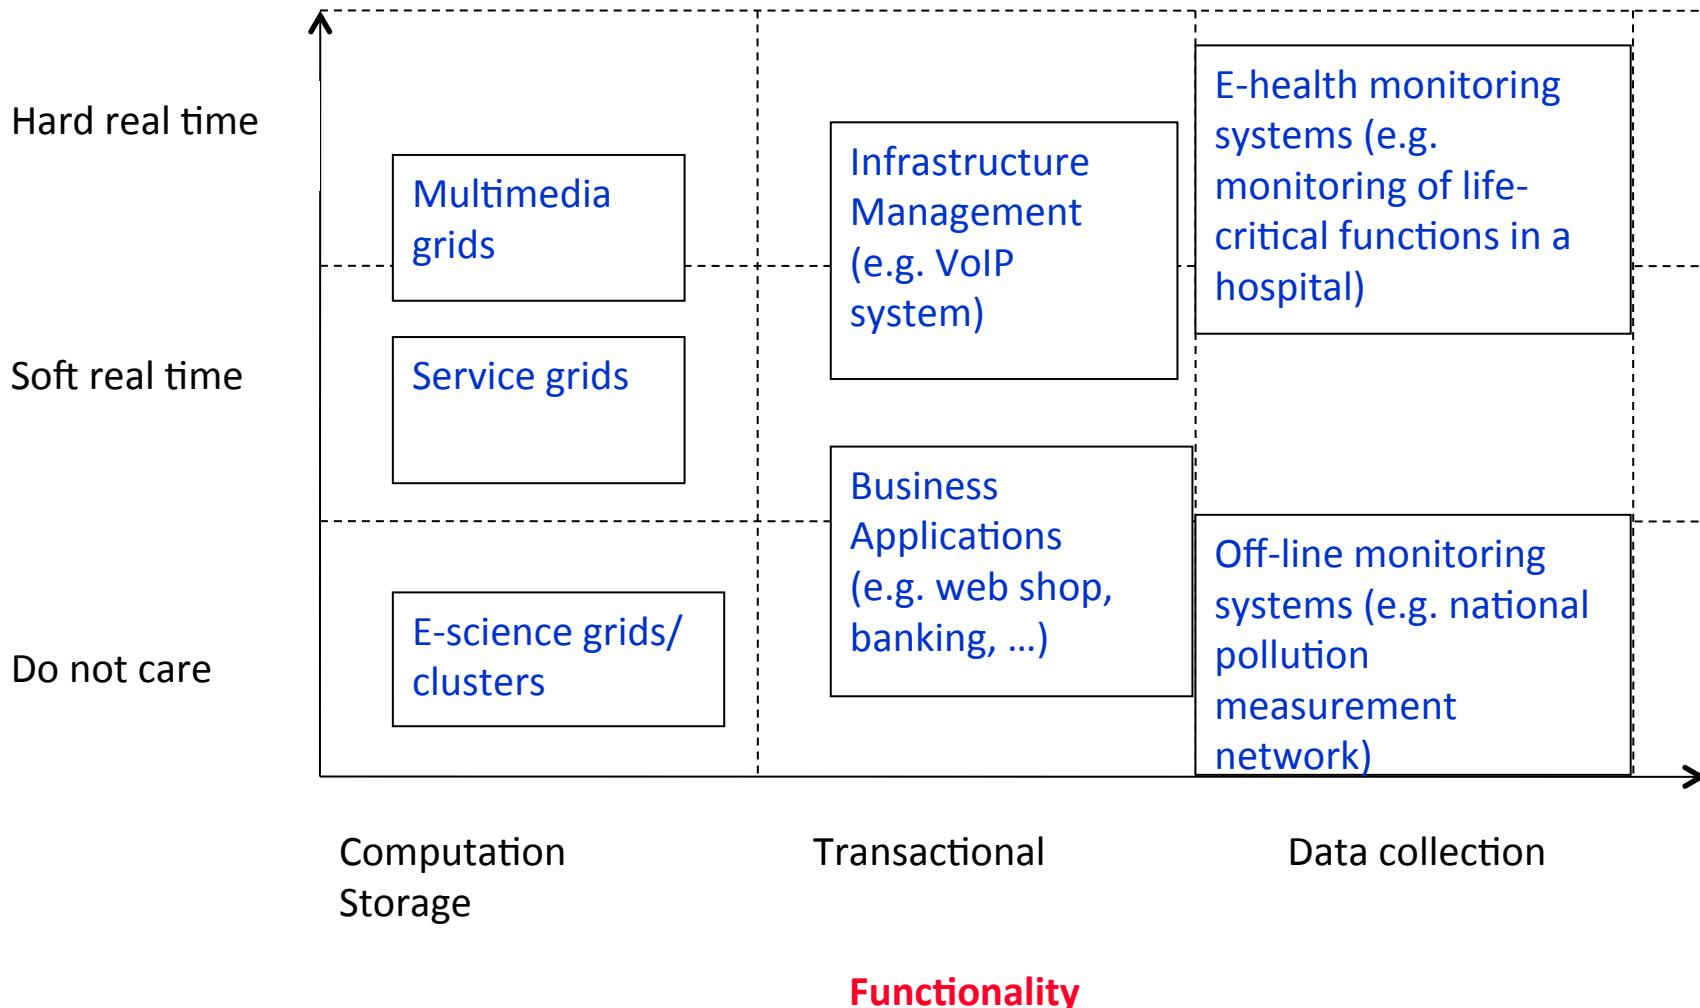
## Chapter 1

# Introduction

1. Definitions and Terminology
2. Developing distributed applications
3. Architecture
4. Design: middleware and services
- 5. Classes of distributed systems**
6. Important architectures and platforms

...

### Responsiveness



## Chapter 1

# Distributed Software: Introduction

1. Definitions and Terminology
2. Developing distributed applications
3. Architecture
4. Design: middleware and services
5. Classes of distributed systems
- 6. Important architectures and platforms**
  1. Thin versus thick ?
  2. Grid and cluster systems
  3. Transactional systems
  4. Integration
7. Scalability and high availability

...

## Thin client

- + cheaper and many devices involved -> overall cheaper
- + cheaper to update (controlled by server), and less updates needed
- some functions more natural to support on client
- permanent network connection needed

## Technologies

### - web browser

- + ubiquitous and uniform client
- + no network blockages (firewalls) for HTTP (port 80)
- + huge user base
- + browser functionality can be extended

### - remote desktop clients

(cf. Athena) “screen scrapers”

## Cluster middleware

- job scheduling
- cluster monitoring
- user administration
- matchmaking
- job submission

Important technologies: Oracle Fusion, Mosix

## Grid middleware

- meta-scheduling
- abstraction of heterogeneity
- information service

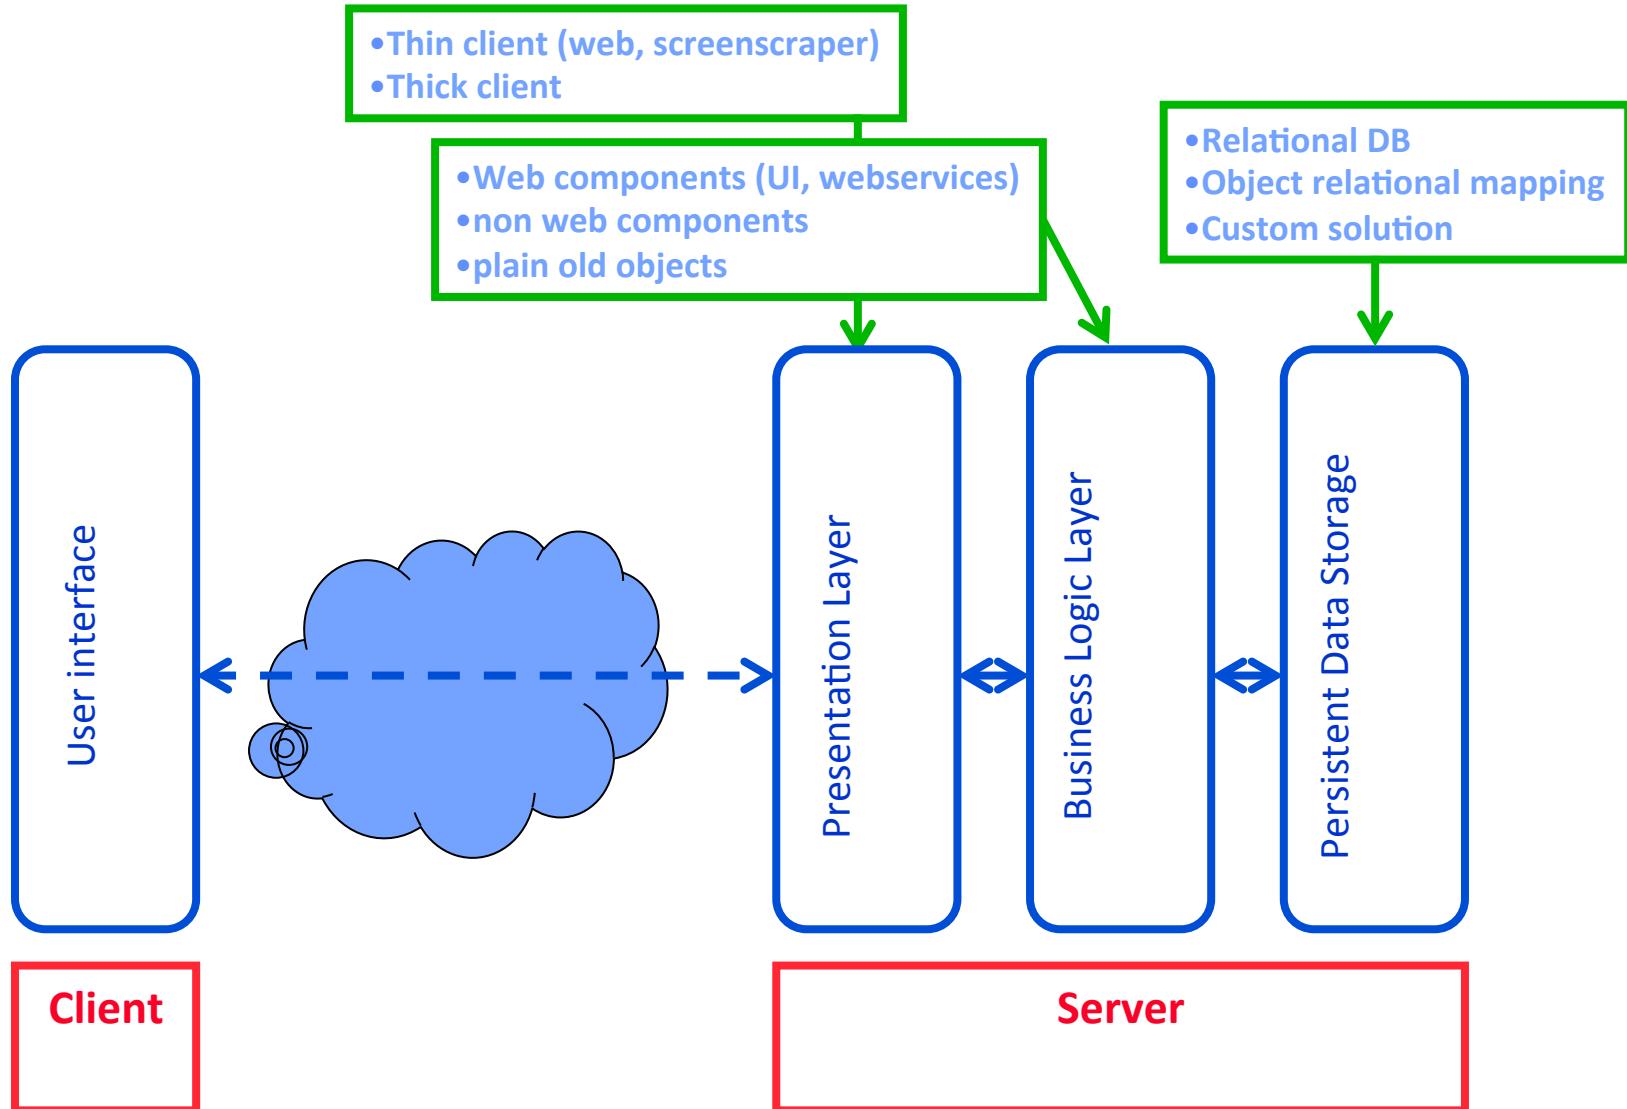
Important technologies :

Globus (and descendants), gLite (used in EGEE(+))

interacting jobs : Message Passing Interface (MPI)



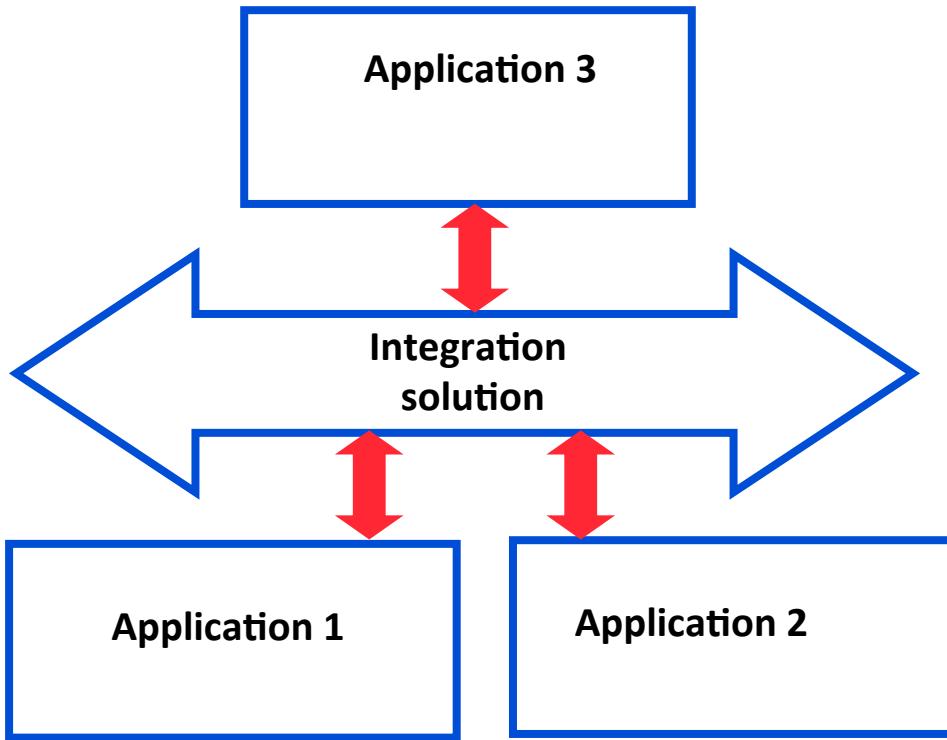




## 3-tier business architecture

layered architecture

rationale : separation of concerns



## Integration solutions :

- Enterprise service bus (SOAP based)
- Message Oriented Middleware

## Chapter 1

# Distributed Software: Introduction

...

**3. Architecture**

**4. Design: middleware and services**

**5. Classes of distributed systems**

**6. Important architectures and platforms**

1. Thin versus thick ?
2. Grid and cluster systems
3. Transactional systems
4. Integration

**7. Scalability and high availability**

**8. The Java technology family**

...

= the system ability to handle gracefully a growing amount of requests

two types of scalability:

- horizontal scalability
- vertical scalability

adding more resources to a single computational node  
improve existing code to handle more requests

### Techniques:

- adding memory or disks, or more powerful CPU
- improving the I/O and concurrency models
  - using non-blocking and/or asynchronous I/O,
  - using the most appropriate thread model (i.e single, pool, per-connection),
  - using synchronous or asynchronous invocations.
- Distribution of processes over available processor cores

adding more computational nodes to distributed system

### Techniques:

- Master/slave techniques
- Partitioning of data on the available nodes  
(e.g. Distributed Hash Tables)
- NoSQL-based systems in particular allow for horizontal scaling

- the property of a distributed system to handle hardware and software failures

### Techniques:

- Redundancy: software components are available on several locations,
- Failover : when a software component or resource fails, automatic failover to the redundant components
- Replicas: multiple copies of the same data are available.

Very often the source of performance problems!

A distributed computer system can NOT simultaneously provide all three of the following guarantees:

- Consistency: all read/write operations must result in a global consistent state,
- Availability: all requests on non-failed components must get a response,
- Partition Tolerance: the system continues to operate when nodes are not able to communicate with each other.

### In practice:

- the Availability and Consistency guarantees are the most important
- Partition Tolerance or effects of network failures are often neglected.

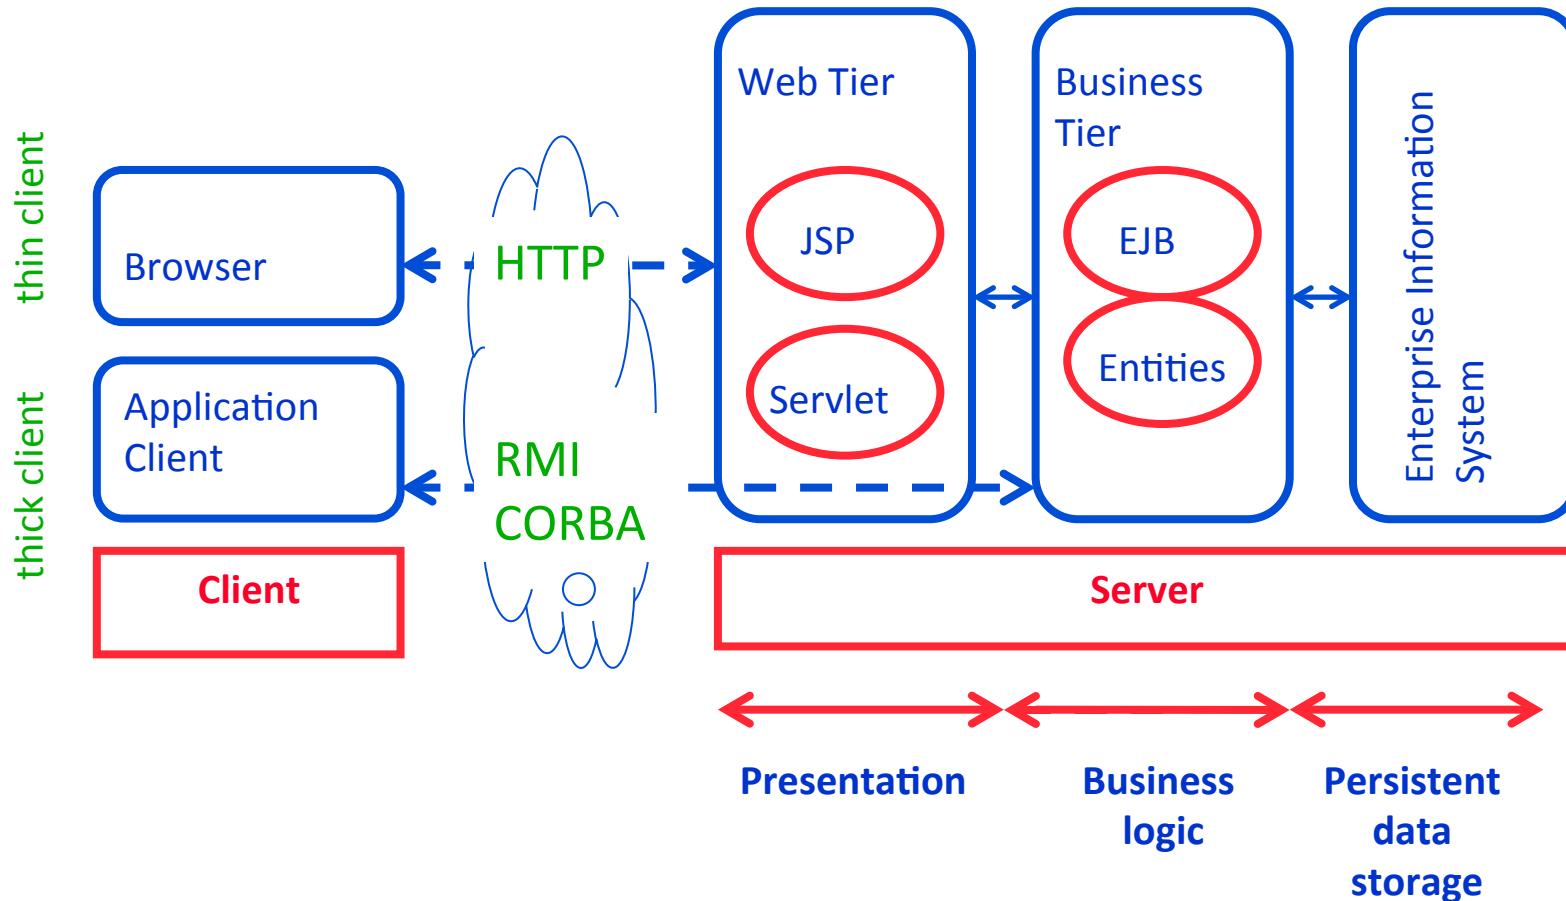
## Chapter 1

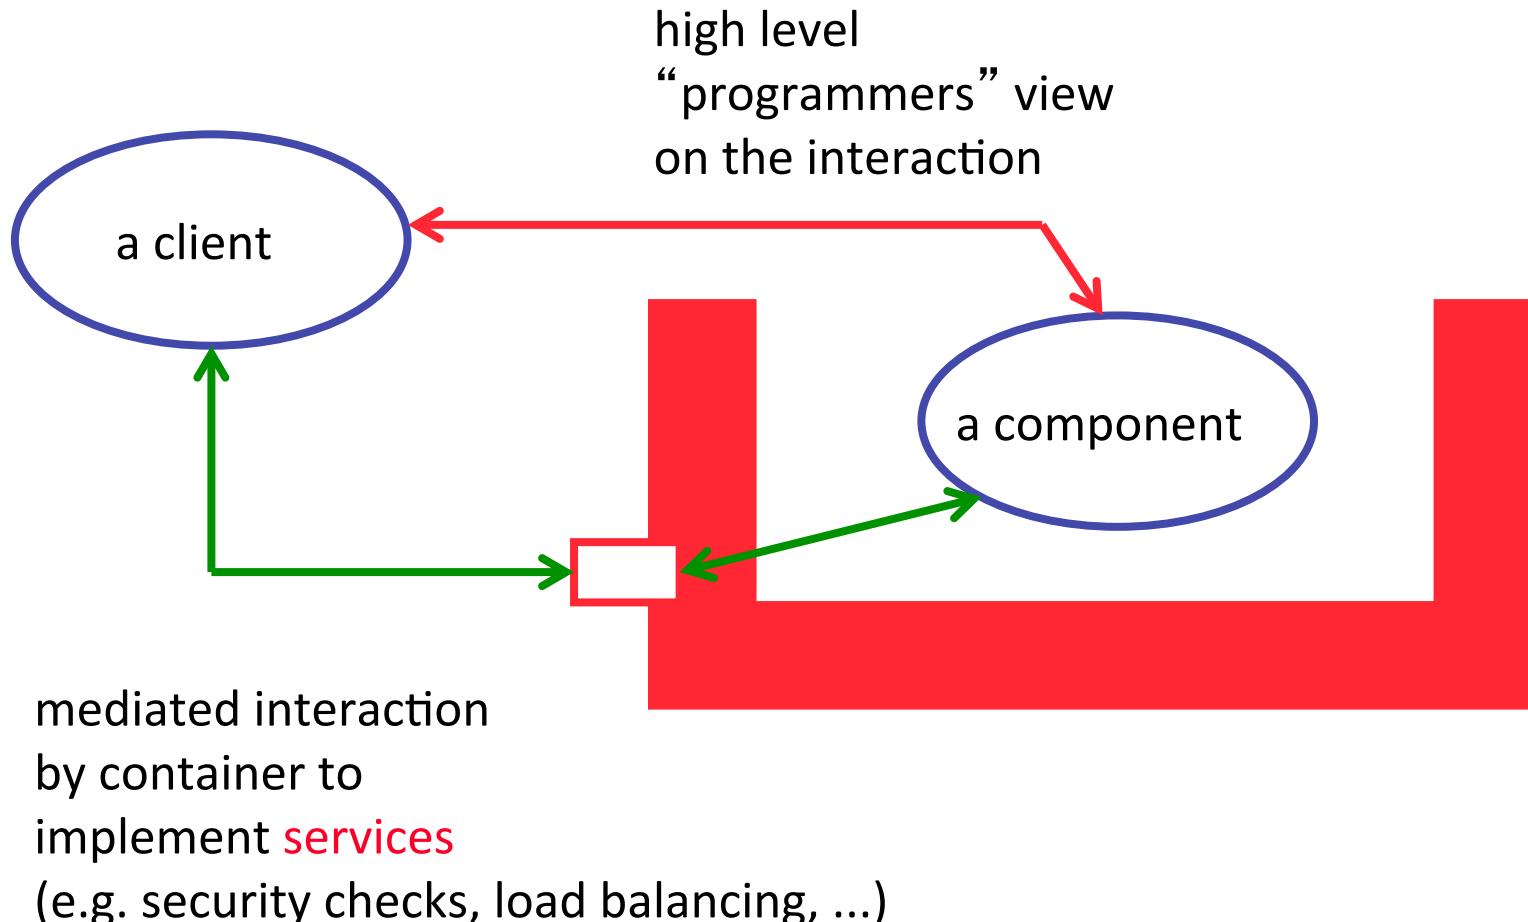
# Distributed Software: Introduction

...

- 4. Design: middleware and services
- 5. Classes of distributed systems
- 6. Important architectures and platforms
- 7. Scalability and high availability
- 8. The Java technology family**
  - 1. General architecture
  - 2. Client tier
  - 3. Web tier
  - 4. Business tier
  - 5. Web services

...





## Thin client

- Container = **browser**
- runs HTTP protocol
  - services
    - install plug-ins
    - run security checks
    - sandboxing

## Thick client

- Container = **appclient**
- runs RMI/CORBA/... protocol
  - middleware services:
    - naming and lookup
    - security checking
    - ...

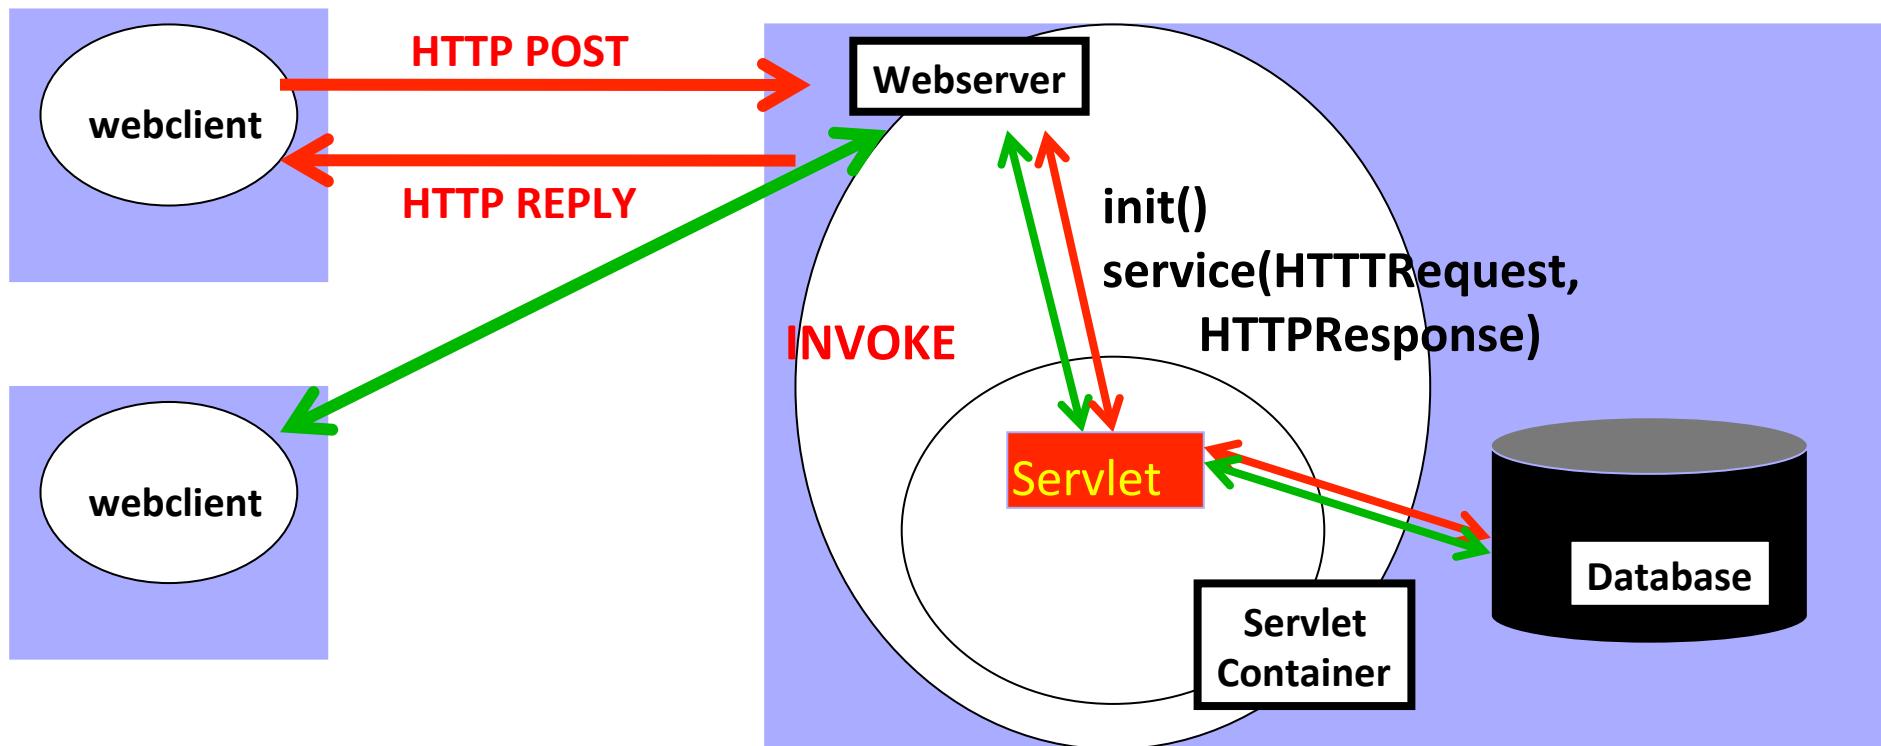
- = server side plug-in
- = base technology for all Java web tier components

## Basic operation

1. browser sends request to URL referring to a Servlet instance  
HTTP request contains parameters, HTTP headers, ...
2. Servlet gets HttpServletRequest object from web server
3. Servlet analyzes request
4. Servlet prints HTML output to HttpServletResponse object
5. HttpServletResponse is sent to webserver
6. Webserver sends HTML to client

## Container services

- Servlet life cycle management
- Servlet pooling
- Session tracking
- Providing access to other resources (e.g. a database)
- Providing access to other components (other Servlets, EJBs, ...)



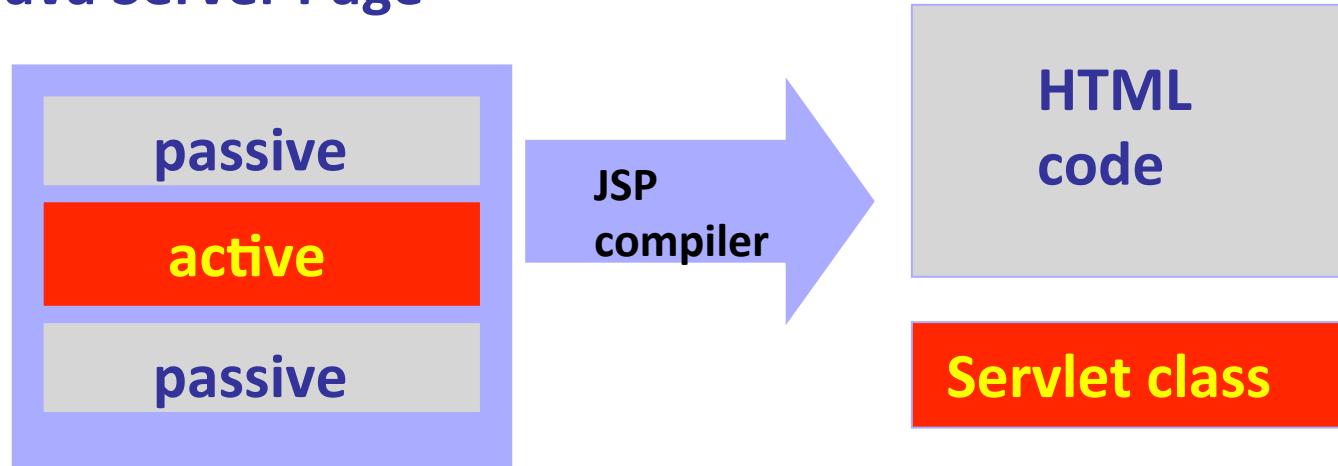
## Servlet disadvantage

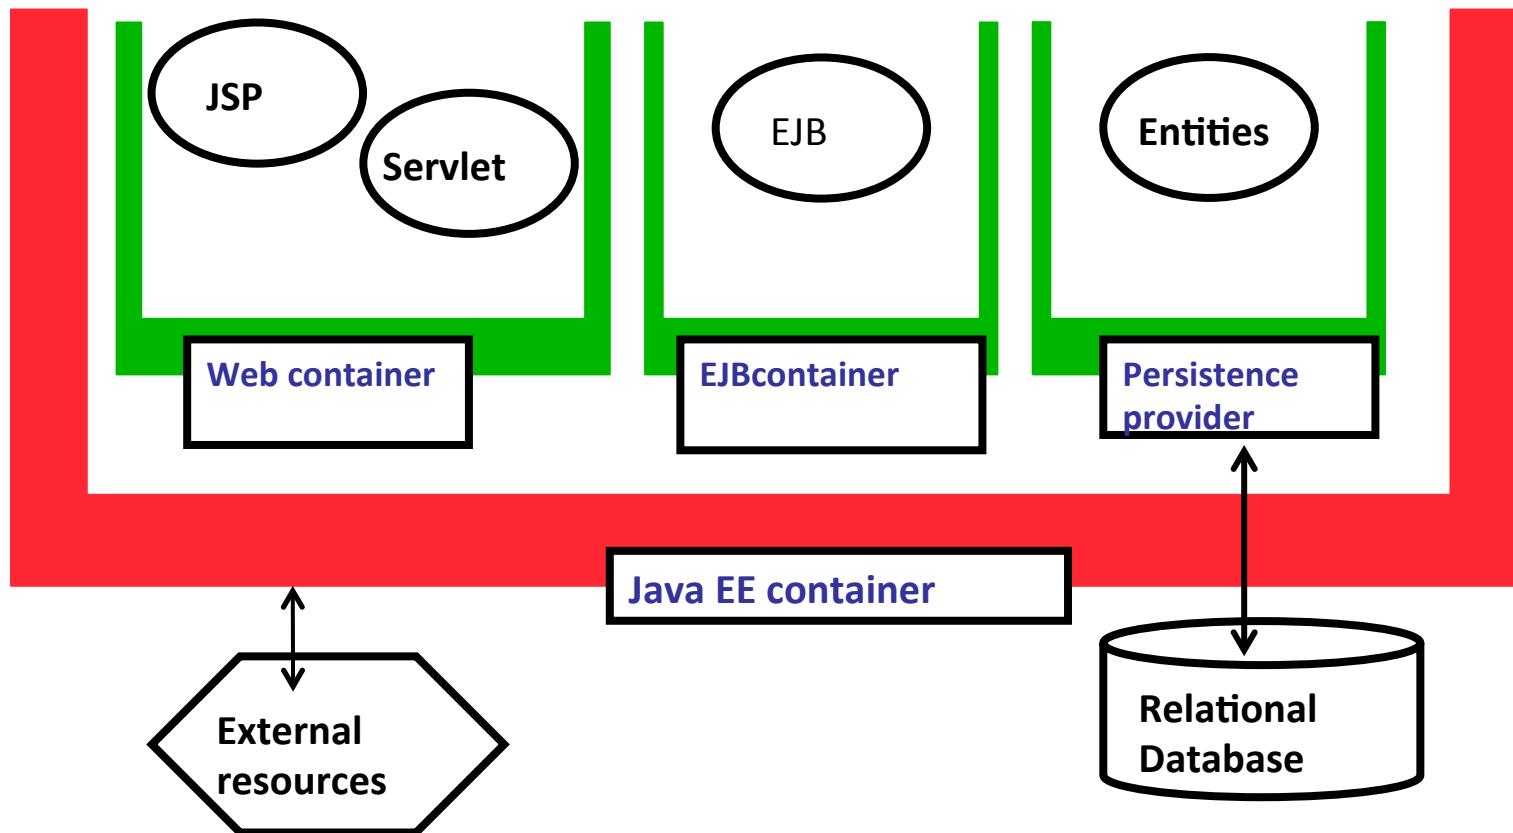
- pages contain static + dynamic data
- static data created dynamically -> not optimal

## Solution

- mix static and dynamic code in single webpage
- compile webpage to extract servlets and static portions
- configure webserver to
  - automatically invoke servlets to create dynamic content
  - merge dynamic and static content to response to client

## JSP : Java Server Page





## EJB container services

- component life cycling management,
- transaction processing,
- security handling,
- persistence
- remotability
- timer
- state management
- resource pooling
- messaging

## Persistence provider services

- store/retrieve from DB
- optimization through caching

## Session Beans

- session related object
- always associated to one single client at most
- types
  - stateful** : “conversational state”
  - stateless**
  - singleton**

synchronous

## Message Beans

- asynchronous message handling
- connected to JMS compliant message queue
- new since J2EE 1.3

asynchronous

## Entities

- “real life” object
- mostly associated to “row in database”
- persistent
- NEW since EJB3.0
- [**replace (very) complex EntityBeans**]
- NOT managed by EJB container

## What ?

= “interface” technology

specifies how to access remote objects through the web

typically : SOAP (Simple Object Access Protocol) or

REST (Representational State Transfer) over HTTP

## Pro – con ?

+ language and platform neutral

+ often uses HTTP as underlying protocol (no network discontinuities)

- performance (large messages, complex parsing)

- only stateless invocations

## JEE and web services

Pojo's and session beans can be exposed as webservice

## Chapter 1

# Distributed Software: Introduction

...

4. Design: middleware and services
5. Classes of distributed systems
6. Important architectures and platforms
7. Scalability and high availability
8. The Java technology family
9. The .NET technology family

...

### Very similar functions as in JEE

#### Web tier

- web service (SOAP / REST) heavily used
- ASP is counterpart for JSP
- .NET framework for web tier development : ASP.NET

#### Business tier

- COM+ component model ported to .NET
- “.NET enterprise services”
  - resource pooling
  - eventing
  - security
  - transaction support
  - synchronization

## **Chapter 1**

# **Distributed Software: Introduction**

...

- 4. Design: middleware and services**
- 5. Classes of distributed systems**
- 6. Important architectures and platforms**
- 7. Scalability and high availability**
- 8. The Java technology family**
- 9. The .NET technology family**
- 10. This course**

- Chapter 1: Introduction
- Chapter 2: Middleware
- Chapter 3: Enterprise Applications
- Chapter 4: Global State and Time
- Chapter 5: Coordination
- Chapter 6: P2P systems
- Chapter 7: Cloud Computing
- Chapter 8: Resource Allocation